LABEL NOISE TECHNIQUES TO FAIRNESS IN MACHINE LEARNING

Ygor de Mello Canalli

Rio de Janeiro
Maio de 2024

LABEL NOISE TECHNIQUES TO FAIRNESS IN MACHINE LEARNING

Ygor de Mello Canalli

TESE SUBMETIDA AO CORPO DOCENTE DO INSTITUTO ALBERTO LUIZ COIMBRA DE PÓS-GRADUAÇÃO E PESQUISA DE ENGENHARIA DA UNIVERSIDADE FEDERAL DO RIO DE JANEIRO COMO PARTE DOS REQUISITOS NECESSÁRIOS PARA A OBTENÇÃO DO GRAU DE DOUTOR EM CIÊNCIAS EM ENGENHARIA DE SISTEMAS E COMPUTAÇÃO.

Orientadores: Geraldo Zimbrão da Silva
Filipe Braida do Carmo

Aprovada por: Prof. Leandro Guimarães Marques Alvim
Prof. Nome do Segundo Examinador Sobrenome
Prof. Nome do Terceiro Examinador Sobrenome
Prof. Nome do Quarto Examinador Sobrenome
Prof. Nome do Quinto Examinador Sobrenome

RIO DE JANEIRO, RJ – BRASIL
MAIO DE 2024

A alguém cujo valor é digno
desta dedicatória.

Agradecimentos

Gostaria de agradecer a todos.

Resumo da Tese apresentada à COPPE/UFRJ como parte dos requisitos necessários para a obtenção do grau de Doutor em Ciências (D.Sc.)

LABEL NOISE TECHNIQUES TO FAIRNESS IN MACHINE LEARNING

Ygor de Mello Canalli

Maio/2024

Orientadores: Geraldo Zimbrão da Silva
               Filipe Braida do Carmo

Programa: Engenharia de Sistemas e Computação

Apresenta-se, nesta tese, uma abordagem ...

Abstract of Thesis presented to COPPE/UFRJ as a partial fulfillment of the requirements for the degree of Doctor of Science (D.Sc.)

THESIS TITLE

Ygor de Mello Canalli

May/2024

Advisors: Geraldo Zimbrão da Silva
            Filipe Braida do Carmo

Department: Systems Engineering and Computer Science

In this work, we present ...

Sumário

Lista de Figuras

Lista de Tabelas

Capítulo 1

Introduction

The issue of fairness in machine learning has recently risen to prominence due to its implications in real-world decision-making systems [1, 2]. Addressing biases and discrimination is a relevant frontier in decision-making systems, as equitable outcomes across various demographic groups is both an ethical imperative and often a legal requirement. Though fairness is a multifaceted concept, it has been deeply examined within the context of machine learning. The literature presents a variety of fairness definitions, drawing concepts from political philosophy and computational techniques [2, 3]. Choosing an equitable machine learning model requires the selection of a fitting definition of fairness, tailored to the specific problem at hand. Many such definitions can be precisely articulated, allowing models to be evaluated based on their predictions.

One inherent challenge in fair machine learning is the balance between fairness and accuracy. Efforts to mitigate unfairness often compromise the model's predictive performance, a trade-off that has been well documented [1, 3]. Predictors that are less biased against marginalized groups may deviate from the true class, resulting in sub-optimal performance. Also, introducing fairness considerations adds constraints to the model, further complicating the optimization process [4].

In light of these challenges, we introduce the Fair Transition Loss, a novel approach to fair classification. This method estimates the influence of historical and societal biases on outcome probabilities for distinct groups within dataset. For instance, individuals from marginalized groups might have lower chances of favorable outcomes compared to their counterparts from privileged groups. Such disparate probabilities can be represented by transition matrices. Drawing inspiration from label noise robustness, we incorporate these transition matrices information into the loss function to promote fairness. The proposed method has some hyperarameters, chosen by a Multi-Objective Optimization approach combining both fairness and model performance with a linear smooth objective. This objective is defined in such a way that it is possible to use this approach to optimize a variety of fairness and

performance metrics.

The primary contribution of this study is the conceptualization of the Fair Transition Loss, a novel loss function influenced by label noise methodologies. In benchmark tests across common fair classification tasks, our empirical results demonstrate that this method consistently outperforms many leading in-processing fair classification techniques in a variety of scenarios. To the best of our knowledge, this work is the first of its kind to apply label noise techniques directly within the model to mitigate unfairness.

The remainder of this paper is structured as follows: Section **??** delineates related works encompassing fairness definitions, metrics and methods, multi-objective optimization in fair machine learning, and classifications in the presence of label noise. We also spotlight studies bridging fairness and label noise. In Section 3.1, we describe the Fair Transition Loss and its underlying principles. Section 4.1 details our experimental methodology and Section 4.2 discuss results for common fair classification tasks. In Section 5 we present conclusions drawn from our study and insights to research directions.

Capítulo 2

Bibliographic revision

2.1 Fairness, Accountability, and Transparency in Machine Learning

The field of Machine Learning (ML) has experienced significant growth and is increasingly applied in various societal domains such as healthcare, finance, and criminal justice. This growth raises important ethical and operational concerns, particularly regarding the principles of Fairness, Accountability, and Transparency (FAT). As ML algorithms increasingly influence a wide array of societal domains, including criminal justice, healthcare, finance, and employment, the imperative to ensure these systems are designed and implemented responsibly has become paramount. This section aims to delineate the significance, scope, and prevailing challenges associated with integrating FAT principles into ML, providing a foundation for the subsequent discussion.

Fairness in ML concerns the equitable and just treatment of all individuals, particularly those from historically marginalized or disadvantaged groups. It seeks to ensure that ML algorithms do not perpetuate existing biases or create new forms of discrimination. However, the multifaceted nature of fairness, encompassing various definitions and metrics, poses substantial challenges in operationalizing it within algorithmic frameworks. Further in this section we will explore these complexities, examining different conceptions of fairness and the inherent trade-offs they entail.

Accountability in ML pertains to the obligation of designers, developers, and deployers of ML systems to be answerable for the outcomes of these systems. It involves establishing mechanisms that allow for the tracing of decisions back to the entities responsible for the deployment of the ML algorithms. Accountability also encompasses the adherence to ethical standards, legal requirements, and societal norms. This discussion frequently involves mechanisms and practices that can enhance accountability in ML, like auditing, documentation, and regulatory

3

compliance.

Transparency, the third pillar, refers to the clarity and openness with which ML systems operate. It involves the ability of stakeholders, including end-users, regulators, and the broader public, to understand how ML systems make decisions. Transparency is crucial for building trust, facilitating informed consent, and enabling the scrutiny necessary to identify and rectify biases. However, achieving transparency, particularly with complex models, presents its own set of technical and ethical challenges. This research topic includes isses as the trade-off between explainability and model performance, and discussing emerging approaches to enhance interpretability without sacrificing effectiveness.

The triad of Fairness, Accountability, and Transparency (FAT) forms the cornerstone of ethical Artificial Intelligence (AI). These principles are pivotal in ensuring that AI systems are developed and deployed in a manner that respects human rights, promotes social well-being, and maintains public trust. While accountability ensures that entities behind AI systems can be held responsible for their outcomes, transparency allows stakeholders to understayfoster environments where AI systems can be scrutinized, understood, and corrected, thereby aligning their functionality with societal norms and values.

However, the present text will concentrate predominantly on the aspect of fairness Fairness is not only crucial for the development of just and equitable technological solutions but also imperative for maintaining the legitimacy and acceptability of AI systems in diverse societal contexts. In delving into the multifaceted dimensions of fairness, this text aims to unpack the theoretical underpinnings, practical challenges, and potential pathways to achieving fairer AI systems, thereby contributing to the broader discourse on ethical AI.

2.1.1 Sources of unfairness

lorem ipsum

2.1.2 Fairness definitions and metrics

This section aims to present some widely used definitions and metrics of fairness, as described by [5] and summarized by [1] and [3], providing a comprehensive overview for understanding and navigating the multifaceted dimensions of fairness in ML systems. Initially, the discourse will delve into general considerations and intuitive aspects of fairness, setting the stage for a deeper understanding. This preliminary discussion is crucial as it lays the groundwork for grasping the nuanced nature of fairness notions within the context of ML. Following this, we will transition into formal definitions, where we will dissect and explain those metrics and concepts.

Even before this discussion, it's crucial to understand that no single fairness definition universally applies to all scenarios. The choice of a particular fairness definition and metric should be informed by ethical considerations grounded in the social context in which the model would be deployed [6]. Selecting a fairness definition is not a purely technical matter, as it inevitably entails ethical and social considerations that should not be neglected [7]. Building fair machine learning models requires an interdisciplinary approach that engages all stakeholders, including specially those who are typically marginalized or underrepresented [8].

A prevalent classification within fairness literature differentiates fairness notions into group metrics and individual metrics. Group Fairness Metrics hinge on the principle that statistical measures — such as error rates, precision, and recall — ought to be equitably distributed across groups demarcated by sensitive attributes like race, gender, or age. The core premise of these metrics is that fairness is actualized when an algorithm exhibits consistent performance across diverse demographic segments.

Demographic Parity, for example, mandates uniformity in the rate of positive algorithmic outcomes across different groups, a standard that remains agnostic to the underlying base rates within each population segment. On the other hand, Equal Opportunity and Equalized Odds introduce a nuance to this conversation by tethering fairness to the true condition of outcomes. Specifically, Equal Opportunity seeks to harmonize the true positive rates across groups, while Equalized Odds extends this alignment to both true positive and false positive rates. This refinement underscores a crucial differentiation within fairness metrics: some are predicated solely on predicted values (such as Demographic Parity), while others derive from the comprehensive landscape of the confusion matrix, incorporating true conditions (as seen in Equal Opportunity and Equalized Odds).

Individual Fairness Metrics, in contrast, introduce a more granular perspective to fairness, advocating that similar individuals should be treated similarly by the ML system. This approach diverges from group-level considerations, focusing instead on ensuring that the algorithm's treatment is consistent for individuals who are alike in relevant aspects, barring their membership in different demographic categories. Individual fairness seeks to ensure a personalized sense of justice, where the algorithmic outcomes are solely reflective of pertinent attributes rather than biased by irrelevant factors associated with sensitive attributes. This concept champions the notion that fairness extends beyond group identities to recognize and respect the uniqueness of individual experiences and qualifications.

One example of Individual Fairness Metric is the notion of Counterfactual Fairness, which introduce a causal reasoning framework into the fairness discourse. These metrics are based on the concept that a decision is fair towards an indivi-

dual if the same decision would have been made in a counterfactual world where the individual belonged to a different demographic group but all other characteristics remained constant. This approach hinges on causal models that specify how sensitive attributes affect other features and the outcome. Counterfactual fairness aims to address the individual-level biases that group fairness metrics might overlook, providing a nuanced approach that considers the hypothetical scenarios of individuals belonging to different demographic categories. By employing counterfactual analysis, one can assess whether the disparities in ML predictions stem from legitimate factors or unjust biases.

These definitions include demographic parity, equal opportunity, predictive equality, and equalized odds. Demographic parity requires that the likelihood of a positive outcome is the same across different groups, irrespective of their sensitive attributes. Equal opportunity extends this concept to the true positive rate, ensuring that individuals from different groups have an equal chance of being correctly classified as positive. Predictive equality, on the other hand, focuses on the true negative rate, ensuring that individuals from different groups have an equal chance of being correctly classified as negative. Equalized odds combines the principles of equal opportunity and predictive equality, ensuring that both true positive and true negative rates are equal across different groups.

Definition 1 (Statistical Parity). *The likelihood of a positive, i.e. favorable, outcome should be the same in every group of the sensitive attribute [9, 10]. A binary predictor $\hat{Y}$ satisfies Statistical Parity (a.k.a. Demographic Parity) if $P(\hat{Y}|A = 0) = P(\hat{Y}|A = 1)$, where A is a protected attribute.*

Demographic Parity does not depend on true class Y, only on prediction $\hat{Y}$. We can measure Demographic Parity (Definition 1) for a protected attribute A as the absolute difference between $P(\hat{Y}|A = 0)$ and $P(\hat{Y}|A = 1)$. According Demographic Parity, the predictor is considered fairer when this metric is lower.

Definition 2 (Equal Opportunity). *The probability of a person in a positive class being assigned to a positive, i.e. favorable, outcome should be the same in every group of the sensitive attribute [11]. A binary predictor $\hat{Y}$ satisfies Equal Opportunity if $P(\hat{Y}|A = 0, Y = 1) = P(\hat{Y}|A = 1, Y = 1)$, where Y is true class and A is a protected attribute.*

Definition 2 claims that protected and unprotected, i.e. privileged, groups should have equal true positive rates. Mathematically, a classifier with equal true positive rates will also have equal false negative rates. Like in Demographic Parity, we can measure Equal Opportunity as an absolute difference between protected and privileged group rates. A variation of Equal Opportunity is Predictive Equality (Definition 3).

Definition 3 (Predictive Equality). *The probability of a person in a negative class being assigned to a negative outcome should be the same in every group of the sensitive attribute. A binary predictor $\hat{Y}$ satisfies Predictive Equality if $P(\hat{Y}|A = 0, Y = 0) = P(\hat{Y}|A = 1, Y = 0)$, where Y is true class and A is a protected attribute.*

Equalized Odds is a combination of the principles from Definition 2 and Definition 3, i.e., protected and unprotected groups should have equal true positive and true negative rates, therefore equal false positive and false negative rates. We can measure Equalized Odds as the average between Equal Opportunity and Predictive Equality.

Definition 4 (Equalized Odds). *Both probabilities of the person in a positive class being assigned to a positive outcome and of a person in a negative class being assigned to a negative outcome should be the same in every group of the sensitive attribute [11]. A binary predictor $\hat{Y}$ satisfies Equalized Odds (a.k.a. Average Odds Difference) if $P(\hat{Y}|A = 0, Y) = P(\hat{Y}|A = 1, Y)$, where Y is true class and A is a protected attribute.*

2.1.3   The impossibility theorem

The Impossibility Theorem, as elucidated by Kleinberg, Mullainathan, and Raghavan, articulates a fundamental challenge in the domain of algorithmic fairness: the concurrent satisfaction of distinct fairness metrics is inherently unfeasible under certain conditions. This theorem, also referred to as the Incompatibility of Fairness Criteria, delineates the intrinsic conflicts arising amongst three prevalent fairness constructs: Group Calibration, Positive Class Balance, and Negative Class Balance.

Group Calibration, or predictive parity, mandates that for any given score produced by the algorithm, individuals across different groups should exhibit equivalent probabilities of actualizing the predicted outcome, be it loan default or recidivism. Positive Class Balance necessitates that the average score assigned by the algorithm to the positive class (e.g., defaulters or recidivists) remains consistent across various demographic groups. Analogously, Negative Class Balance demands uniformity in the average score assigned to the negative class (e.g., non-defaulters or non-recidivists) across these groups.

The crux of the Impossibility Theorem is the assertion that, in contexts where base rates (the actual occurrence rates of the event in question) diverge between groups, it is mathematically untenable for an algorithm to simultaneously fulfill all three aforementioned fairness criteria. This revelation does not denote a shortfall of the algorithm per se but rather highlights the statistical conundrums intrinsic to comparing disparate groups.

The ramifications of the Impossibility Theorem extend deeply into the fabric of machine learning and societal decision-making. It posits that in the algorithmic design phase, stakeholders are compelled to engage in critical deliberations to prioritize among conflicting fairness criteria, given the impracticality of achieving universal fairness under every metric. This necessity for prioritization underscores the imperative roles of transparency and accountability, where the selected fairness criteria and the justification for their preference should be transparently communicated and justified to all stakeholders involved.

Moreover, the theorem advocates for a departure from universalist approaches towards algorithmic fairness, promoting instead a context-sensitive paradigm. In this paradigm, the decision regarding which fairness criterion to prioritize should be informed by the unique societal, legal, and ethical contexts pertinent to each specific application of machine learning. Such a nuanced approach to fairness is paramount for the development of machine learning systems that not only exhibit technical robustness but also align with overarching societal values and ethical principles.

2.1.4   Fair classification

In this section, we review pertinent literature on fair machine learning, placing a particular emphasis on in-processing techniques. We also delve into classification methodologies that operate in the presence of label noise and discuss multi-objective optimization within the context of fair machine learning. While our research does not directly tackle fairness problems in the presence of label noise, we highlight relevant works that, akin to ours, bridge the domains of fairness and noise in machine learning research.

Fairness intervention methods can be classified into three categories based on the stage at which they occur, as proposed by [1] and [6]: Pre-processing methods intervene before learning, modifying the data to reduce existing biases; In-processing methods intervene during learning by modifying the objective functions or imposing constraints to the model in order to mitigate discriminatory effects; Post-processing methods affects predictions produced by the model after learning to change possibly unfair outcomes. In this work, we incorporate information about disparities among social groups in the dataset into our model by modifying the loss function through the use of a transition matrix. This fairness intervention is thus classified as an in-processing technique.

Other relevant in-processing strategies for fair classification include Naive Bayes approaches for discrimination-free classification [12], Fairness Through Awareness Framework [9], Fairness-Aware Classifier with Prejudice Remover Regularizer [13], α-discriminatory empirical risk minimizer [14], Disparate Impact and Disparate Mis-

treatment frameworks for margin-based classifiers [4, 15], Weak Agnostic Learning to Auditing Subgroup Fairness [16, 17], One-Network Adversarial Fairness [18], FairGan$^+$ [19], Monte Carlo policy gradient method [20], Fairness-accuracy Pareto [21], and Pareto front stochastic multi-gradient [22] based in original stochastic multi-gradient [23] to Multi-Objective Optimization and the hybrid Adaptive Priority Reweighing approach [24].

Recently, special attention has been given in fair machine learning research topics like addressing multiple sensitive attributes or multiple classes [25, 26], loss balancing techniques [27, 28], where the objective is to balance the loss across different groups instead of predictive metrics, adversarial approaches [29–34] and the privacy concerns involving fairness under federated learning settings [35, 36]. Another relevant research topic in fair machine learning is learning under censored data [37–40], which we will discuss in section 2.2.3.

2.1.5 Fairness and multi-objective optimization

A model that substantially decreases model performance to reduce unfairness may not be a viable option, as low performance could harm all groups affected by the model's decisions, including protected groups. Similarly, a model projected to be a fair alternative that keeps performance almost intact, but with little or even no gain in fairness, is not practically relevant. It is possible that fine-tuning this trade-off could result in a fairer solution that achieves better performance than traditional methods, but this is not the case for most practical problems. Achieving this balance is one of the most challenging tasks in fair machine learning.

In this context, an interesting approach is to deal with fair machine learning as a Multi-Objective Optimization (MOO) problem, where predictive performance and fairness metric are the objectives, which could be defined according Equation 2.1, where λ is a parameter configuration in the space Λ, $\rho : \Lambda \mapsto [0,1]$ is a model performance metric and $\varphi : \Lambda \mapsto [0,1]$ a fairness metric. The set of all optimal solutions is called Pareto front, where one objective cannot be improved without sacrificing another. In this setting there is no single λ^* optimal solution, but a set of solutions forming a Pareto front [41].

$$\max \ (\rho(\lambda), \varphi(\lambda)) \tag{2.1}$$
$$\text{subject to } \lambda \in \Lambda$$

One of the most frequent approach to deal with MOO problems like these is to combine the multiple function outputs to a single scalar, which is called scalarization. Therefore, we could describe a general scalarization setup to Equation 2.1 according

Equation 2.2. The effectiveness of this approach is that is also possible to use single objective optimization techniques to tackle the MOO optimization problem. In this scenario a relevant issue is to select a scalarization setup capable of promote a proper trade-off of all the objectives thorough the optimization process given the optimization method.

$$\arg\max_{\lambda \in \Lambda} G(\lambda) = (\rho(\lambda), \varphi(\lambda)) \tag{2.2}$$

The fairness-accuracy Pareto front is formally described in [21], which demonstrate that many existing fairness methods are performing a linear scalarization scheme and argues that it has several limitations in recovering Pareto optimal solutions. Instead, authors proposes a Chebyshev scalarization scheme, that is theoretically superior than linear scheme. A characterization of the accuracy-fairness trade-off as a Pareto front can be found in [22]. Also, [23] proposes a stochastic multi-gradient based in original stochastic multi-gradient to Multi-Objective Optimization.

Another remarkable use of MOO in Fair Machine Learning is to perform a Fair Hyperparameter Optimization, which can offer a model agnostic approach with flexibility to apply in multiple machine learning pipelines. A time-efficient Baysian Optimization approach can be found in [42], combining scalarization techniques with the bandit-inspired Hyperband [43] algorithm to Hyperparameter Optimization in context of fairness.

A general objective function to be used with some popular off-the-shelf hyperparameters optimization techniques combining model performance and fairness in a flexible setting can be found in [44]. The authors argues that in fairness context the Pareto front is most often convex, thus proposes a simple scalarizing function that could be applied to reduce G to a single scalar with weighed l_p-norm. Also, they argue that [45] demonstrate the the use of l_p-norms with a high p value leads to slower convergence. Thus, the optimization metric $g(\lambda) = ||G(\lambda)||_1$ is optimized according Equation 2.3, where α is the relative importance of predictive performance and fairness and λ is a parameter configuration in the space Λ. In experiments, α is fixed at 0.5, giving same importance to both objectives.

$$G(\lambda) = \alpha \cdot \rho(\lambda) + (1 - \alpha) \cdot \varphi(\lambda) \tag{2.3}$$

A Multi-objective SVMOptimizer with Dataset Constraints is proposed by [46], where the objective is to minimize multiple objectives on real-world datasets, such as misclassification error and positive prediction at specific rate to some population. A custom reinforcement learning algorithm directly modeling performance and fairness as objectives is proposed by [20]. Authors proposes using as reward function

the difference between model performance (Area Under the ROC Curve) and three different fairness metrics (Statistical Parity, Equal Opportunity and Equalized Odds), each one with its respective importance coefficient. In experimental setups only one of those coefficient are different from zero. Thus, the optimized metric could be written as $G(\lambda) = \rho(\lambda) - \alpha \cdot \varphi(\lambda)$, where α is the relative importance of fairness.

2.2   Classification in the presence of label noise

The presence of noise in data can substantially decrease model performance in classification problems. Noise can be defined as non systematic errors that obscures the relationship between features of an instance and its class [47–49]. Two types of noise are found in literature, in features (or attributes) and in labels (or classes). Feature noise affects observed values, e.g. by adding a small Gaussian noise on each feature during measurement. Likewise, label noise change the observed label assigned to an instance, e.g. by randomly inverting labels in a binary classification problem. Although feature noise could affect model performance, label noise is potentially more harmful, since we frequently have many features and only one label. Note that in label noise only the observed label of an instance is affected, its true class remains the same.

2.2.1   Label noise taxonomy

The label noise taxonomy considers three types of noise: Noisy Completely at Random, Noisy at Random, and Noisy Not at Random [47]. Figure 2.1 presents the statistical dependency between features X, class Y, observed label $\tilde{Y}$ and the occurrence of error E, i.e. $E = 1$ when $Y \neq \tilde{Y}$. The simplest type is Noisy Completely at Random, where the occurrence of error E not depend on X and Y, e.g. randomly flipping labels on a binary classification problem. In Noisy at Random, the occurrence of error E depends only on Y, e.g. randomly flipping labels on binary classification with different rates for positives and negatives classes. Noisy Not at Random considers the occurrence of error E depending on both Y and X, e.g. flipping labels on binary classification with different rates for each group of instances of a certain feature.

2.2.2   Loss factorization and correction

Many label noise robustness methods can be found on literature, in this work we highlight the *backward* and *forward* loss corrections, proposed by [50] using concepts

Figura 2.1: Noise taxonomy from a statistical perspective. (a) completely random noise (NCAR), (b) random noise (NAR) and (c) non-random noise (NNAR). The arrows correspond to the statistical dependencies. For clarity, the dependency between X and Y was placed as a dashed arrow.

of loss factorization [51]. Those loss correction techniques considers a NAR label noise, which is described by a transition matrix T such as

$$T_{i,j} = P(\tilde{Y} = y_j | Y = y_i), \qquad (2.4)$$

where $\mathcal{Y} = \{y_1, y_2, \ldots, y_c\}$ is the set of all possible class labels. Transition matrix includes corruption probabilities for every possible label combination, each value represents the probability of one label be corrupted onto another. This matrix is row-stochastic and not necessarily symmetric across the classes.

$$\ell^{\leftarrow}(P(\tilde{Y}|X)) = T^{-1}\ell(P(\tilde{Y}|X)) \qquad (2.5)$$

The backward loss correction is defined by Equation 2.5 to an arbitrary loss function ℓ and a transition matrix T. The backward loss correction involves a linear combination of the loss values for each observed label, using coefficients that depends on the probability that each observed label reflects the true class. Intuitively, we are reweighting the loss according to the noise probabilities of each label using the inverse of T and thus somehow going one step back, reverting the noise effects. This corrected loss is unbiased and can be minimized with any conventional back-propagation algorithm, making it flexible to include within different training techniques and data pipelines.

$$\ell^{\rightarrow}(P(\tilde{Y}|X)) = \ell(T^{\top}P(\tilde{Y}|X)) \qquad (2.6)$$

However, backward correction requires matrix inversion, which may not exist or may lead to numerical instabilities if the transition matrix T is ill-conditioned. Although there is possible solutions to a bad condition number of T, one should consider using the forward correction, a backward variation proposed by [50] to avoid this issue, as defined in Equation 2.6. While backward acts on the loss itself, forward corrects model predictions. Forward correction does not have the same

theoretical guarantees as backward, but offers a label noise robustness, ensuring that the learned model is the minimize over the clean distribution without the need of matrix inversion.

2.2.3   Fairness in the presence of noise

Some recent works deal with fairness problems in the presence of noise. For example, the sensitive attribute available could be noisy, which could distort the effects of fairness intervention. In this context, [52] uses noise-rate estimators from the label noise literature to change a fairness model. Also, [53] proposes a framework for assessing how assumptions on the noise across groups affect the predictive bias properties in risk assessment models. Furthermore, [54] considers the consequences of naively relying on noisy protected group labels while proposing two new optimization approaches with sensitive attribute noise robustness. A denoised version of the selection problem to deal with noisy sensitive attributes is proposed in [55]. Lastly, [56] proposes an optimization framework for classification in the presence of noisy protected attributes.

There is also the perspective of dealing with the proxy features divergence or covariance. A theoretical approach to this issue identifying potential sources of errors can be found in [57]. The problem of measuring group fairness in ranking based on divergence with proxy features is investigated by [58]. A framework of fair semi-supervised learning in the pre-processing phase can be found in [59], which includes predicting labels for unlabeled data, a resampling method, and ensemble learning to improve accuracy and decrease discrimination.

Another research direction is considering how fair models perform in the presence of NNAR label noise, where error rates of corruption depend both on the label class and the membership of a protected subgroup. In this scenario [60] addresses the problem of fair classification and [61] provides a general framework for rewriting the classification risk and the fairness metric in terms of noisy data and thereby building robust classifiers. In [62] a study about the presence of noise in the protected attribute can be found.

Furthermore, many recent works deals with fairness under semi-supervised settings considering censored data, that is, for some individuals the class label is not available due censorship [37–40]. In this scenario, the main approach is to use some technique to estimate the missing data instead of removing the instance from training data. This is closely related to the previous problems of fair learning under noisy data. In censored fairness problems noise can be interpreted as a kind of censorship, as the original data affected by noise is not available.

Bias and noise are two related phenomena, both corrupt data affecting models

trained with this data. For example, if noise disproportionately affects different groups this potentially produces unfairness in models that uses this data in training [60]. For example, we could have positive true class ($Y = 1$) flipped into negative labels ($\tilde{Y} = 0$) more frequently in the protected group ($A = 1$) than in privileged group ($A = 0$). Simultaneously, the negative class ($Y = 0$) could be more frequently flipped into positives observed labels ($\tilde{Y} = 1$) within privileged/unprotected group ($A = 0$). This scenario could lead to a undetected higher false negative rate to protected group and higher false positive rate to privileged group. In this case the Noisy Not at Random data would be a source of negative social bias.

In [1] a non-exhaustive list of bias types was presented, including *Historical Bias*, *Representation Bias*, *Sampling Bias*, *Omitted Variable Bias*, among others. In the scenario described above, the incorrect measurement of the true class resulted in a different observed label ($Y \neq \tilde{Y}$), which could be classified as a *Measurement Bias*. Similarly, a Noisy Not at Random data could lead to a *Population Bias*, where the characteristics of the population represented in the data differ from those of the original target population.

It can be challenging to distinguish between label noise and bias in certain scenarios, specially when noise disproportionately affects different social groups. Although there is some overlapping, they are distinct phenomena. Label noise is a stochastic process that is considered independent and unintentional [47], whereas bias is rooted in historical and social issues and could be intentional. Furthermore, even noise-free data, correctly represented by observed features and labels, may be unfair since the social phenomena that produce this data could be biased against some groups.

Prior studies in the realm of fairness have largely concentrated on understanding how noisy or censored data affects fair learning and on mitigating these effects. In contrast, our approach is inspired by label noise techniques, but with a distinct goal: not merely to analyze or mitigate the impact of noise or censorship, but to directly address and reduce unfairness itself.

Capítulo 3

Proposal

3.1 Fair Transition Loss

We propose a novel fair classification method inspired by techniques used for classification in the presence of label noise. By leveraging the features of certain label noise methods that redistribute probabilities for unbalanced noise across classes, our approach re-weights prediction probabilities to reduce disparities in favorable and unfavorable outcomes across social groups.

Whereas forward loss correction [50] uses a transition matrix with corruption probabilities for every label combination in the case of NAR, fair classification problems are more related to NNAR. While forward loss correction uses a transition matrix with corruption probabilities for each label combination, as in the case of NAR, fair classification problems align more with NNAR scenarios. In NNAR, the probability of corruption depends not only on the true class but also on features, analogous to how bias in fairness problems is directed against certain groups. Here our correction does not revert a random label corruption from the true class, but a potentially unfair prediction. While noise label techniques, like forward [50], aims to correct the prediction targeting a unknown true class using the available noisy label, analogously the proposed technique focus on correcting predictions chasing the unknown fair class using the available unfair label. Despite those are distinct phenomena, the corrections works the same way, adjusting the probabilities of predictions produced by a machine learning model during the training.

Thus, our proposal is a prediction probability loss reweighting technique that accounts different rates to each group of the sensitive feature, instead of using the same correction to every individual. A correction method that incorporates different probabilities for protected and unprotected groups could be more effective in mitigating bias during the learning phase. Specifically, we want a forward-based correction that takes into account not only one transition matrix, but a different matrix

to each group of sensitive features. In this scenario, each group of sensitive feature have its own correction, with its own rates for each class combination. Ideally, if we can find an appropriate transition matrix that describes the bias to each group in a specific problem, we can apply a correction that attenuates those negative effects by reweighting model's predictions in the learning process.

Next, we formally present Fair Transition Loss. For purpose of clarity we follow the same structure available at [50], with the pertinent changes to our scope. The Fairness Transition Matrix T_a is defined with some abuse of notation to the group $A = a$ of the sensitive feature as

$$T_{a,i,j} = P(\tilde{Y} = y_j | Y = y_i, A = a), \tag{3.1}$$

where label space $\mathcal{Y} = \{y_1, y_2, \ldots, y_c\}$, c the number of classes, $Y = y_i$ is the unknown fair class and $\tilde{Y} = y_j$ is the available and possibly unfair label. Here, $T_{a,i,j}$ is the probability of the fair class $Y = y_i$ being unfairly labeled as $\tilde{Y} = y_j$ to an individual of the group $A = a$ due negative social bias. Therefore, suppose that there is an inversible link function $\psi : \Delta^{c-1} \to \mathbb{R}^c$, where $\Delta^{c-1} \subset [0,1]^c$ is the c-simplex, the simplex in a c-dimensional space. Thus, a composite loss function, denoted by $\ell_\psi : \mathcal{Y} \times \mathbb{R}^c \to \mathbb{R}$ if it can be written as a decomposition of ψ^{-1}, that is,

$$\ell_\psi(Y, h(X)) = \ell(Y, \psi^{-1}(h(X)), \tag{3.2}$$

where $h : \mathcal{X} \to \mathbb{R}^c$ is a standard artificial neural network with multiple layers using activation functions, and $h(X)$ is the output o this neural network to a given input X. For example, to cross entropy loss function the softmax is the inverse link function. Proper loss functions are those that can be directly used to estimate class probabilities. The minimizer of a proper composite loss has the particular form of the link function applied to the conditional class probabilities $P(Y|X)$. Adding a new conditioning to this formulation, to an individual from group $A = a$ we have

$$\arg\min_h \mathbb{E}_{X,Y} \ell_\psi(Y, h(X|A = a)) = \psi(P(Y|X, A = a)). \tag{3.3}$$

Fair Transition Loss consists in correcting model's predictions with the same technique as forward, but taking into account the sensitive attribute value when choosing the transition matrix. In Theorem 1 the Fair Transition Loss is formally defined, with a guarantee about its minimizers.

Theorem 1. *Suppose that the Fairness Transition Matrix T_a for a given sensitive attribute $A = a$ is non-singular. Given a proper composite loss ℓ_ψ, define the Fair*

Transition Loss as

$$\text{FTL}_\psi(h(X|A = a)) = \ell(T_a^\top \psi^{-1}(h(X|A = a))).$$

Then, the minimizer of the corrected loss under the unfair distribution is the same as the minimizer of the original loss under the fair distribution:

$$\arg\min_h \mathbb{E}_{X,\tilde{Y}} \text{FTL}\, \psi(Y, h(X|A = a)) = \arg\min_h \mathbb{E}_{X,Y} \ell_\psi(Y, h(X|A = a)).$$

Demonstração. First notice that:

$$\begin{aligned}
\text{FTL}_\psi(Y, h(X|A = a)) &= \ell(Y, T_a^\top \psi^{-1}(h(X|A = a))) \\
&= \ell_\phi(Y, h(X|A = a)),
\end{aligned} \tag{3.4}$$

where we denote $\phi^{-1} = \psi^{-1} \circ T_a^\top$. Equivalently, $\phi = (T_a^{-1})^\top \circ \psi$ is invertible by composition of invertible functions, its domain is Δ^{c-1} as of ψ and its codomain is $\mathbb{R}^c$. The last loss in Equation 3.4 is proper composite with link ϕ. Finally, from Equation 3.3, the loss minimizer over the unfair distribution is

$$\begin{aligned}
\arg\min_h \mathbb{E}_{X,\tilde{Y}} \ell_\phi(Y, h(X|A = a)) &= \phi(P(\tilde{Y}|X, A = a)) & (3.5) \\
&= \psi((T_a^{-1})^\top) P(\tilde{Y}|X, A = a) & (3.6) \\
&= \psi(P(Y|X, A = a)), & (3.7)
\end{aligned}$$

that proves the Theorem by Equation 3.3 once again. □

Considering a common scenario with only two groups in sensitive attributes (protected and privileged), we can correct the model's predictions using two different fair transition matrices. One with rates applied while learning instances from the protected group, and the other with rates applied while learning instances from the privileged group. Formally, to the sensitive feature $A \in \{0, 1\}$, let T_0 the transition matrix associated with privileged/unprotected group ($A = 0$) and T_1 with the protected group ($A = 1$), FTL can be computed as

$$\text{FTL}(P(\tilde{Y}|X)) = (1 - A) \cdot \ell(T_0^\top P(\tilde{Y}|X)) + A \cdot \ell(T_1^\top P(\tilde{Y}|X)), \tag{3.8}$$

which in a standard batch learning, consists in alternating the transition matrix applied according instance's sensitive attribute.

Furthermore, to a common binary classification problem, where there is a positive (favorable) class and a negative (unfavorable) class, and two groups from sensitive

feature (protected and privileged), we have two 2×2 transition matrices. Intuitively we are choosing rates to increase or decrease the probability of each group to be classified with the positive or negative prediction. We name those rates associated with increasing the probability to achieve the positive outcome as *promotion* rate, and those associated with increasing the probability to receive the negative outcome as *demotion* rate. As the transition matrix is row-stochastic, we can describe T_0 and T_1 as

$$T_0 = \begin{bmatrix} 1 - d_0 & d_0 \\ p_0 & 1 - p_0 \end{bmatrix}, \; T_1 = \begin{bmatrix} 1 - d_1 & d_1 \\ p_1 & 1 - p_1 \end{bmatrix}, \tag{3.9}$$

where d_0 is the privileged demotion rate, p_0 the privileged promotion rate, the d_1 protected demotion rate, and p_1 the protected promotion rate. With an appropriate combination of d_0, p_0, d_1, p_1 we can define a transition matrix pair that should be able to reweight model's predictions with FTL to achieve fairer results with a reasonable model performance. The central problem in our methodology thus relies in choosing these rates, which can be seen as an hyperparameter optimization problem.

Our hyperparameter optimization problem consists in finding an optimal trade-off between fairness and performance, which can be described as a MOO problem, as defined in Equation 2.1. Here, the hyperparameter configuration is $\lambda = (d_0, p_0, d_1, p_1)$. Since the transition matrix is row stochastic these parameters are sufficient to define T_0 and T_1. We want to maximize model performance $\rho(\lambda)$ and minimize fairness metric $\varphi(\lambda)$.

$$G(\lambda) = \rho(\lambda) - \varphi(\lambda). \tag{3.10}$$

Following some MOO approaches to fair machine learning, we will use a linear scalarization setup to define the optimization metric [20, 42]. As we yet have four hyperparameter to fine-tune, and in [44] the relative importance α is fixed at 0.5, we choose a simple and intuitive objective function in Equation 3.10 to maximize without the parameter α, i.e., giving same importance to fairness and performance. In Equation 3.10 we establish a simple objective to optimize, but one might need to consider a different formulation depending on the specific problem at hand.

Capítulo 4

Experimental evaluation

4.1   Experimental setup

In this section, we detail the experimental setup employed to benchmark our model against relevant in-processing fair classification models found in standard fairness toolkits, namely, Prejudice Remover [13], Adversarial Debiasing [63], and Gerry Fair Classifier [16]. We use the implementation of these methods from AI Fairness 360 toolkit [64]. The baseline is a Standard MLP using two hidden layers with 100 hidden units each, *ReLU* activation function, batch size of 64, 50 epochs early stopped at 3 epochs without improvement [65] and softmax in output, trained with ADAM optimizer [66] with learning rate at $3e-4$. The only difference between baseline MLP and Fair Transition Loss MLP is that baseline uses standard Binary Cross Entropy Loss. The Gerry Fair Classifier implementation uses the False Negative Rate as its fairness definition and in Adversarial Debiasing classifier the hidden size is 100 units. Additionally, we compare the Fair Transition Loss within the Adaptive Priority Reweighting [24], a promising fairness promoting technique focused on improving generalization, which outperformed many recent methods such as [67], [68], and [69].

Our methodology consists of two phases: hyperparameter tuning and testing. In the hyperparameter tuning phase we perform a Bandit-Based pruning approach using HyperBand [70] with Tree-structured Parzen Estimator Sampler (TPE) [71] over 100 trials. Those techniques deliver better solutions to multi-objective hyperparameter optimization in the same number of trials than conventional approaches like Grid Search and Random Search [72]. At each trial fitness function is evaluated by performing a complete training and validation, where both model performance and fairness metrics are assessed. The fitness function is computed based on the objective defined in Equation 3.10. This same experimental procedure can be adapted to utilize other hyperparameter tuning algorithms such as FairRandom Search, Fair TPE, and Fairband [44].

Once the best hyperparameters are selected, we proceed to the testing phase, where a new training is conducted using those optimal hyperparameters. After this training, we evaluate the model's performance on a separate test set that was not used during the hyperparameter tuning phase, which are reported. This complete tuning-training-testing described is repeated 15 times with dataset re-sampling then we proceed to comparison. Here the re-sampling consists in shuffling the whole dataset before splitting, which is better described further in this section.

As the objective defined in Equation 3.10 can be achieved with different performance and fairness metrics, we compare the proposed method with other relevant in-processing techniques from literature in different optimization scenarios. For predictive performance, we evaluate not only Accuracy (Acc.) but also the Mathews Correlation Coefficient (MCC), which has advantages over F1 score and Accuracy in binary classification evaluation [73], where 1 means a perfect prediction according true class, -1 a complete inversion and 0 an average random outcome. As fairness metric we consider Statistical Parity (Stat. Parity, Definition 1), Equal Opportunity (Eq. Opp., Definition 2) and Equalized Odds (Eq. Odds, Definition 4). Thus we have the following optimization scenarios: MCC and Statistical Parity; MCC and Equal Opportunity; MCC and Equalized Odds; Accuracy and Statistical Parity; Accuracy and Equal Opportunity; Accuracy and Equalized Odds.

Tabela 4.1: Hyperparameters search ranges or options of each method.

Method	Parameter	Range/options
Standard MLP (baseline)	dropout	[0.0, 0.2]
Prejudice Remover [13]	η	[0.0, 50.0]
Adversarial Debasing [63]	α	[0.0, 1.0]
Gerry Fair Classifier [16]	C	[0.0, 20.0]
	γ	{0.1, 0.01, 0.001}
Fair Transition Loss	d_0, p_0, d_1, p_1	[0.0, 1.0]
	dropout	[0.0, 0.2]

Table 4.1 presents the methods hyperparameters along with their corresponding search ranges or options. While each method may possess a varying number of hyperparameters and range sizes, all are optimized under the same conditions and number of configurations to guarantee a balanced comparison. In Figure 4.1 we present a brief sensibility analysis on the fitness functions with different fairness and performance values over those six optimization scenarios listed before. Here we perform a complete hyperparameter tuning with HyperBand and TPE over 100 trials using the baseline model (Standard MLP) with the Adult Income dataset optimizing the hyperparamers reported on Table 4.1. This sensibility analysis aims better present the linear objective function behavior within performance and fairness

metrics evaluated in this study.

Figura 4.1: Sensibility analysis on optimized fitness functions within different performance and fairness metrics. Results from complete hyperparameter tuning through 100 trials with baseline model over the Adult Income dataset.

Each plot in Figure 4.1 illustrates the fitness value corresponding to specific performance and fairness metrics. The color scheme in the plots represents the fitness values, with higher values in red and lower values in blue. On the x-axis, we have the fairness metric, where a lower value is preferable. The y-axis represents the performance metric, with higher values being preferred. The color gradient in these plots demonstrates the linear relationship between fitness and variations in the corresponding performance and fairness metrics. It is important to note that the scale of the fairness metric is significantly smaller than that of the performance metric. However, it is sufficiently to act as a penalization. The general fitness function in the scenarios described is capable of producing results with reduced bias while maintaining similar performance levels. Although each metric combination has different scales, each hyperparameter tuning experiment uses only one metric combination at a time, ensuring consistency in the optimization process.

On these plots we used the fitness and performance levels obtained through the TPE sampler with HyperBand pruning during the hyperparameter tuning phase using the baseline model, as previously described. In this setting, the solution (i.e., the combination of hyperparameters) that yields the best fitness value is selected for a new complete training phase. This involves assessing metrics on test data not used in the previous phase. To ensure robust evaluation, the dataset is reshuffled, re-split into train, validation and test segments, and this entire process is repeated over 15 iterations.

To properly compare this set of 15 results of each method, we conduct an Almost

Stochastic Order (ASO) test [74], which is a significance test suitable for comparing complex machine learning models with various hyperparameters. The ASO test involves evaluating a set of metrics through multiple samplings of a Collection of Statistics (in this case assessed in test phase using random resampling) to compare one method against another. The $ASO(A, B)$ function yields a value in the range $[0, 1]$, given two methods A and B. If $ASO(A, B)$ is lesser than 0.5, we can reject the null hypothesis and conclude that method A outperforms method B in the given task. That is, method A produces stochastically larger values than method B for a given metric. The lower the $ASO(A, B)$ value, the stronger the evidence that A is superior to B in that particular task, which can be interpreted as a confidence interval. Therefore, we perform pairwise comparisons between all methods for each optimization scenario outlined previously and for each dataset.

Our experiments uses common datasets used in Fair Classification problems, namely *Adult Income* [75], *German Credit* [76], *Bank Marketing* [77], and *COMPAS Recidivism* [78]. We use the dataset readers available in the AI Fairness 360 toolkit [64] with its standard configurations. Instances with missing data are removed.

Tabela 4.2: Dataset details used in this work, including performance and fairness metrics assessed to a standard classifier without tuning, and the maximum correlation between sensitive feature and the other features.

Dataset	Adult Income	Bank Marketing	COMPAS Recidivsm	German Credit
# Features	102	57	401	58
# Instances	45222	30488	6167	1000
Sensitive Attribute	sex	age	race	sex
Positives	24.78%	12.66%	54.45%	70.00%
Negatives	75.22%	87.34%	45.55%	30.00%
Privileged	67.50%	97.17%	34.05%	69.00%
Unprivileged	32.50%	2.83%	65.95%	31.00%
Accuracy	0.846	0.906	0.358	0.685
MCC	0.572	0.553	-0.275	0.000
Stat. Parity.	0.192	0.106	0.172	0.074
Equal Opportunity	0.094	0.145	0.120	0.043
Equalized Odds	0.092	0.094	0.163	0.122
Maximum Correlation	0.527	0.364	0.826	0.593

Table 4.2 present dataset details used in this work, including the number of features before pre-processing, the count of valid instances, the proportion of positive and negative labels, the sensitive feature considered in experiments, the proportion of privileged and unprivileged groups within the corresponding sensitive feature, reference performance and fairness metrics using a standard Random Forest Classifier with 1000 classifiers without tuning, and the maximum correlation coefficient between the sensitive feature and the other features. The maximum correlation is useful to assess whether it is possible to use another feature as proxy to the sensitive

feature, which is commonly referred as redlining effect [79].

The *Adult Income* dataset presents a classification task to predict whether an individual earns more than $50,000$ per year. This dataset consists of $48,842$ instances sourced from the U.S. 1994 Census database. The sensitive attribute used in this task is sex, with the male group considered privileged and the female group protected (unprivileged). In the *German Credit* dataset, the task consists of classifying $1,000$ individuals described by a set of attributes as good or bad credit risks. Similar to the *Adult Income* dataset, here we use sex as the sensitive attribute, with the male group considered privileged and the female group protected. The *Bank Marketing* classification task aims to predict whether $45,211$ clients will subscribe to a term deposit after direct marketing campaigns (phone calls) by a Portuguese banking institution. In this case, the sensitive feature is age, where individuals under the age of 25 are considered unprivileged, while those aged 25 and older are considered privileged. The *COMPAS* dataset presents around $80,000$ criminal records from the Broward County Clerk's Office. The task here is to predict whether a defendant will recidivate in the next two years. The sensitive feature in this case is race, with Caucasians as the privileged group and non-Caucasians (Black and Hispanic) as unprivileged.

For all datasets, the data preparation process is the same, one-hot encoding for categorical features and standardize the numerical features. We perform a random split, with 80% allocated for the hyperparameter tuning phase and the remaining 20% reserved for evaluating metrics in the test phase. Within the hyperparameter tuning phase, this corresponding fraction of data is further randomly split, with 80% assigned to training and 20% to validation. The validation set allows us to assess metrics and compute the fitness function for each hyperparameter configuration. In datasets where there is originally some kind of split (e.g., train set and test set in separate files), all available data is merged and then shuffled to produce new splits at each run.

4.2   Results and discussion

This section summarizes our results, comparing Fair Transition Loss (FTL) with the baseline Standard MLP (MLP) and some relevant fair in-processing methods: Adversarial Debiasing (AD), Prejudice Remover (PR), Gerry Fair Classifier (GFC) and Adaptive Priority Reweighting (APW).

As we have multiple optimization scenarios with different objective functions and datasets, and to each of them multiple runs, we present in Table 4.3 the results of the ASO test described before, which allow us to properly compare each method to FTL. Values under 0.5 (in bold) mean that we can reject the null hypothesis,

Tabela 4.3: Almost Stochastic Order test comparing Fair Transition Loss fitness. Values under 0.5 (in bold) mean that FTL outperforms corresponding method in such optimization scenario.

Fairness/Performance Metric	Dataset	MLP	AD	PR	GFC	APW
Statistical Parity MCC	Adult Income	**0.00**	**0.15**	1.00	**0.00**	1.00
	Bank Marketing	**0.01**	**0.00**	**0.00**	**0.00**	**0.00**
	Compas Recidivism	**0.01**	**0.25**	**0.00**	**0.02**	**0.00**
	German Credit	**0.28**	**0.30**	**0.39**	**0.21**	**0.28**
Equal Opportunity MCC	Adult Income	**0.01**	**0.00**	**0.05**	**0.00**	0.93
	Bank Marketing	0.81	**0.18**	**0.24**	**0.09**	0.77
	Compas Recidivism	**0.00**	1.00	**0.00**	0.66	1.00
	German Credit	1.00	**0.23**	0.84	0.78	0.76
Equalized Odds MCC	Adult Income	**0.03**	**0.28**	**0.42**	**0.00**	**0.00**
	Bank Marketing	**0.46**	**0.18**	**0.12**	**0.02**	**0.18**
	Compas Recidivism	**0.01**	0.58	**0.00**	**0.07**	**0.00**
	German Credit	1.00	**0.07**	1.00	**0.31**	1.00
Statistical Parity Accuracy	Adult Income	**0.01**	**0.26**	**0.32**	**0.00**	0.53
	Bank Marketing	**0.25**	1.00	1.00	0.76	0.82
	Compas Recidivism	**0.00**	1.00	**0.10**	1.00	**0.00**
	German Credit	1.00	**0.26**	1.00	1.00	1.00
Equal Opportunity Accuracy	Adult Income	0.89	0.97	1.00	**0.23**	0.98
	Bank Marketing	1.00	**0.39**	0.81	1.00	1.00
	Compas Recidivism	**0.01**	0.78	**0.00**	**0.10**	1.00
	German Credit	1.00	0.64	1.00	1.00	1.00
Equalized Odds Accuracy	Adult Income	**0.01**	**0.21**	**0.19**	**0.00**	1.00
	Bank Marketing	0.76	**0.40**	0.82	1.00	1.00
	Compas Recidivism	**0.01**	**0.45**	**0.00**	**0.01**	**0.00**
	German Credit	1.00	**0.12**	1.00	1.00	1.00

i.e., FTL produces stochastically larger fitness than method in respective column for a objective and dataset. Lower values indicate stronger evidence. The complete results with mean and standard deviation of fitness, performance and fairness can be found in Appendix 4.3 to a fairness-performance trade-off analysis.

In 69 of 120 comparison scenarios from Almost Stochastic Order test (Table 4.3), it is possible to claim that FTL outperforms its competitor, i.e., FTL produces stochastically higher fitness values. Despite these positive results, one can argue that the proposed technique only adds extra hyperparameters that increase models flexibility to achieve higher fitness values. In other words, are we effectively describing bias in datasets by transition matrices as claimed before? To address this, we showcase various FTL hyperparameter combinations selected during the tuning phase described in Section 4.1, comparing with corresponding dataset information available at Table 4.2. We perform this analysis using *Adult Income* dataset. The corresponding hyperparameters can be found in Table 4.4. Here, high values mean that FTL alters the corresponding probabilities, while values close to zero indicate

minimal interference by the method.

Tabela 4.4: Fair Transition Loss hyperparameters chosen by optimizing different metrics in *Adult Income* dataset.

Objective	d_0 Priv. Dem.	p_0 Priv. Prom.	d_1 Prot. Dem.	p_1 Prot. Prom.
MCC and Stat. Parity	0.056	0.076	0.043	0.878
MCC and Eq. Opp.	0.292	0.455	0.329	0.575
MCC and Eq. Odds	0.037	0.165	0.005	0.432
Acc. and Stat. Parity	0.470	0.110	0.023	0.446
Acc. and Eq. Opp.	0.389	0.326	0.311	0.530
Acc. and Eq. Odds	0.497	0.286	0.228	0.094

When optimizing for MCC and Statistical Parity, there's a notable high value for protected promotion. This value is compatible with the high corresponding fairness metric for this dataset, approximately 0.19 without correction. This enhances the likelihood that an unprivileged instance receives a favorable outcome. Since statistical parity only compare the probability of a positive outcome across groups (ignoring true class) this is enough. The other hyperparameters presents low values. In contrast to the previous case, optimizing Equal Opportunity requires compatible false negative rates. Optimizing this fairness metric within MCC produces the effect of promoting both privileged and protected, although protected with higher values. This produces the effect of reducing false negatives at all, since the method enhances the probability of a positive outcome. This effect is counterbalanced with intermediate demotion rates to both groups through a finetuning to keep MCC. Note that Equal Opportunity values without correction to this dataset is not as high as Statistical Parity. To optimize Equalized Odds within MCC it is necessary to keep both false negatives and false positives comparable across groups, which lead to a less intense intervention when compared to Equal Opportunity. Here remains the high values to protected promotion to achieve fairness.

There is a remarkable difference between hyperparameters found through optimizing Accuracy and MCC. While MCC handles unbalanced classes effectively, Accuracy only measures the probability of correctly predicting an instance. If the dataset is unbalanced it is possible to achieve high Accuracy only by predicting the label of the more frequent class. In this dataset, only about a quarter of the instances are positives, which can lead to more frequent negative outcomes to achieve higher Accuracy. Results optimizing for Accuracy show significantly higher demotion rates compared to those from MCC optimization, both to privileged and protected groups. From this analysis, it's evident that the proposed methodology effectively describes and mitigates bias in a dataset according to a given fairness definition and while keeping targeted performance metric at a reasonable level.

Now we discuss the results to each objective, starting with MCC and Statistical Parity. To this objective Fair Transition Loss consistently outperforms all methods at all classification tasks, except Prejudice Remover (PR) and Adaptive Priority Reweighting (APW) in *Adult Income.* Figure 4.2a presents a box-plot comparison, where we can see that FTL, PR and APW are effectively drawn. While FTL and APW has little bit higher values, PR presents smaller variance. AS PR is a regularized logistic regression, it is a smaller model than FTL, which can explain the also smaller variance.

Figura 4.2: Fitness values optimizing MCC and multiple fairness metrics.

(a) Statistical Parity   (b) Equal Opportunity   (c) Equalized Odds

When comparing the optimization results for MCC and Equal Opportunity, FTL consistently outperforms its counterparts in most scenarios. Here we have only one discrepancy, as APW presents slightly advantage over FTL on COMPAS dataset. Also, on German Credit dataset all methods achieved similar results. Given the small size of this dataset, we theorize that all methods, barring AD, have reached the Pareto front — meaning, any further improvements in fairness would necessitate a proportional sacrifice in performance. This equilibrium between baseline (MLP), FTL, PR, GFC and APW is evident in Figure 4.2b.

The sub-optimal results of AD can likely be attributed to the dataset's limited size; with merely 1000 instances, this dataset might be too small for the an adversarial model like AD train effectively. This pattern of AD underperforming persists across subsequent classification tasks involving this dataset. Interestingly,

also with the exception of AD, we notice that the variance in results for most optimization scenarios is smaller than in other classification tasks. This observation further underscores our Pareto front hypothesis and suggests that the classification task's simplicity may contribute to the reduced variability in outcomes.

When optimizing for MCC and Equalized Odds, we find that the results are consistent, with FTL outperforming its counterparts in most scenarios. Notably, within the *German Credit* dataset, FTL surpasses not just AD but also GFC. Since GFC primarily relies on the False Negative Rate for its fairness definition, it has a natural advantage when optimizing for Equal Opportunity compared to Equalized Odds, which requires maintaining equitable False Positive and False Negative Rates. As observed in the previous comparisons, Figure 4.2c underscores that the baseline, FTL, PR and APW seem to hit the Pareto front for this dataset.

Upon examining FTL's results when optimizing for MCC across all the fairness metrics evaluated, it's evident that FTL consistently delivers superior results compared to its counterparts. Specifically, FTL achieves stochastically higher fitness values in 44 out of the 60 scenarios evaluated. Given the inherent challenges associated with optimizing MCC compared to Accuracy, we attribute FTL's dominance in the MCC optimization to its capability of effectively capturing the bias idiosyncrasies of the dataset and the specified performance and fairness metrics trough transition matrices.

Furthermore, it's noteworthy that both the baseline and PR models substantially underperform across all optimization scenarios in the *COMPAS Recidivism* classification task. This dataset, characterized by its complexity with 401 features, might be at the heart of these subpar results. We theorize that the this lack of performance could be due to an insufficiently large model to navigate such a high-dimensional space, especially when we observe, as indicated in Table 4.2, that the standard performance on this dataset is relatively low.

When turning our attention to results obtained by optimizing Accuracy, we must first reiterate its inherent simplicity as a performance metric compared to MCC. Given its nature, it allows models to attain high values simply by predicting the label of the predominant class. In such circumstances, it is comparatively easier to reach the Pareto front. Even under these conditions, FTL displays commendable competitiveness. Although it achieves stochastically higher fitness values in 25 of the 60 scenarios, this rate is notably less than what we observed when optimizing for MCC. By juxtaposing Figures 4.3a, 4.3b, and 4.3c with Table 4.3, we discern that in scenarios where FTL does not have the upper hand, it still competes closely with its counterparts. This very close results are primarily attributed to multiple methods simultaneously approaching the Pareto front. Likely when optimizing for MCC using Equal Opportunity as fairness metric, APW presents slightly advantage

over FTL on COMPAS dataset.

A particularly straightforward fair classification task emerges when optimizing for Accuracy and Statistical Parity within the *Adult Income* dataset. With this pronounced class imbalance, models can lean into over-predicting the majority class, thereby aligning the probabilities of positive predictions across the groups. This strategy results in a exceptionally low variance across all methodologies. A similar, albeit reduced, effect can be observed within the *German Credit* dataset, as previously highlighted.

Figura 4.3: Fitness values optimizing Accuracy and multiple fairness metrics.

(a) Statistical Parity · (b) Equal Opportunity · (c) Equalized Odds

Fair Transition Loss consistently demonstrates effective bias mitigation, it does so by absorbing the nuances of the dataset and the fairness metric through its transition matrices, resulting in stochastically superior fitness values in a significant number of scenarios. Additionally, the method has the capacity to effectively handle datasets with unbalanced classes when optimizing for metrics like MCC. However, it's important to recognize that Fair Transition Loss requires fine-tuning multiple hyperparameters. We thus consider that this technique is especially beneficial in setups where hyperparameter optimization is an inherent part of the prediction pipeline.

A key concern is about the potential for fairness-promoting techniques to inadvertently shift the burden onto the very group they aim to protect. This arises from the possibility that by imposing additional constraints, the method might uninten-

tionally learn alternative ways to reproduce and even reinforce the negative social biases present in the data, thus harming the individuals it intends to safeguard.

Figura 4.4: Results of false negatives and false positives within groups on protected promotion (p_1) parameter at increasing levels.

To evaluate the capability of FTL to address this risk, we present another experiment, where we adjusted only the protected promotion hyperparameter (p_1, Equation 3.9) during the FTL training, keeping all other FTL hyperparameters at zero and dropout at 0.2. Here we follow the same re-sampling procedure, each experiment is performed over 15 repetitions, shuffling the dataset before splitting. This experiment was conducted using the four datasets previously analyzed, and we reported the following metrics assessed over the training set: protected false negative rate, protected false positive rate, protected accuracy, privileged false negative rate, privileged false positive rate, privileged accuracy, overall false negative rate, overall false positive rate and overall accuracy.

The results, presented in Figure 4.4, show that increasing the protected promotion hyperparameter value leads to a decreased false negative rate for the protected group. Meanwhile, most other monitored metrics tended towards stability to reasonable hyperparameter levels (p_1 under 0.95). An exception is the false positive rate, specially to protected group, which increased as the false negative rate decreased

Tabela A.1: Complete results optimizing MCC and Statistical Parity.

Dataset	Method	Fitness	MCC	Stat. Parity
Adult Income	Adaptive Priority Reweighting	0.499($\pm$0.01)	0.510($\pm$0.01)	0.011($\pm$0.01)
	Fair Transition Loss	0.492($\pm$0.02)	0.512($\pm$0.01)	0.020($\pm$0.01)
	Prejudice Remover	0.491($\pm$0.01)	0.500($\pm$0.01)	0.009($\pm$0.01)
	Adversarial Debiasing	0.478($\pm$0.01)	0.501($\pm$0.02)	0.024($\pm$0.02)
	Standard MLP (baseline)	0.395($\pm$0.01)	0.581($\pm$0.01)	0.185($\pm$0.01)
	Gerry Fair Classifier	0.357($\pm$0.02)	0.512($\pm$0.02)	0.154($\pm$0.03)
Bank Marketing	Fair Transition Loss	0.539($\pm$0.03)	0.579($\pm$0.01)	0.040($\pm$0.03)
	Adversarial Debiasing	0.459($\pm$0.03)	0.505($\pm$0.02)	0.046($\pm$0.02)
	Prejudice Remover	0.454($\pm$0.03)	0.487($\pm$0.02)	0.033($\pm$0.02)
	Adaptive Priority Reweighting	0.441($\pm$0.03)	0.482($\pm$0.02)	0.041($\pm$0.04)
	Standard MLP (baseline)	0.419($\pm$0.04)	0.522($\pm$0.02)	0.102($\pm$0.03)
	Gerry Fair Classifier	0.358($\pm$0.04)	0.428($\pm$0.02)	0.070($\pm$0.03)
COMPAS Recidivism	Fair Transition Loss	0.220($\pm$0.06)	0.276($\pm$0.03)	0.057($\pm$0.05)
	Adversarial Debiasing	0.157($\pm$0.14)	0.322($\pm$0.02)	0.165($\pm$0.14)
	Gerry Fair Classifier	0.141($\pm$0.04)	0.289($\pm$0.06)	0.148($\pm$0.06)
	Prejudice Remover	$-$0.318($\pm$0.05)	$-$0.276($\pm$0.03)	0.042($\pm$0.03)
	Adaptive Priority Reweighting	$-$0.412($\pm$0.35)	0.194($\pm$0.07)	0.606($\pm$0.29)
	Standard MLP (baseline)	$-$0.511($\pm$0.05)	$-$0.299($\pm$0.03)	0.212($\pm$0.04)
German Credit	Fair Transition Loss	0.272($\pm$0.08)	0.354($\pm$0.07)	0.083($\pm$0.04)
	Prejudice Remover	0.234($\pm$0.09)	0.329($\pm$0.05)	0.095($\pm$0.06)
	Standard MLP (baseline)	0.223($\pm$0.10)	0.330($\pm$0.07)	0.107($\pm$0.07)
	Gerry Fair Classifier	0.221($\pm$0.09)	0.291($\pm$0.11)	0.071($\pm$0.06)
	Adaptive Priority Reweighting	0.217($\pm$0.09)	0.321($\pm$0.05)	0.105($\pm$0.06)
	Adversarial Debiasing	0.200($\pm$0.17)	0.368($\pm$0.06)	0.168($\pm$0.15)

to keep accuracy. This pattern was consistent across all evaluated datasets. The primary aim of this experiment was to demonstrate that our proposed technique does not inadvertently penalize the protected group. Rather, the overall impact of increasing the aforementioned parameter is to effectively promote fairness for the protected group without detriment to either the privileged or protected groups. The behavior of the remaining FTL parameters is analogous, necessitating proper fine-tuning to achieve a balanced outcome. This underscores the efficacy of our method in achieving its intended purpose of reducing bias and promoting fairness in the model.

4.3 Complete Results

In this appendix we present complete results with mean fitness (Equation 3.10), performance and fairness across multiple resampling, to a proper trade-off comparison. To each objective function and dataset methods results are ordered from higher to lower fitness mean, with corresponding standard deviation presented between parenthesis. Results corresponding each optimization scenario can be found on tables A.1, A.2, A.3, A.4, A.5, and A.6.

Tabela A.2: Complete results optimizing MCC and Equal Opportunity.

Dataset	Method	Fitness	MCC	Eq. Opp.
Adult Income	Fair Transition Loss	0.523($\pm$0.02)	0.576($\pm$0.02)	0.052($\pm$0.02)
	Prejudice Remover	0.509($\pm$0.05)	0.558($\pm$0.02)	0.049($\pm$0.03)
	Adversarial Debiasing	0.509($\pm$0.03)	0.565($\pm$0.02)	0.056($\pm$0.02)
	Adaptive Priority Reweighting	0.493($\pm$0.01)	0.523($\pm$0.01)	0.030($\pm$0.01)
	Standard MLP (baseline)	0.489($\pm$0.03)	0.576($\pm$0.01)	0.087($\pm$0.03)
	Gerry Fair Classifier	0.434($\pm$0.01)	0.523($\pm$0.01)	0.089($\pm$0.01)
Bank Marketing	Fair Transition Loss	0.485($\pm$0.06)	0.569($\pm$0.01)	0.084($\pm$0.06)
	Standard MLP (baseline)	0.439($\pm$0.03)	0.514($\pm$0.02)	0.075($\pm$0.03)
	Adversarial Debiasing	0.426($\pm$0.06)	0.512($\pm$0.02)	0.086($\pm$0.05)
	Adaptive Priority Reweighting	0.424($\pm$0.04)	0.474($\pm$0.02)	0.050($\pm$0.04)
	Prejudice Remover	0.413($\pm$0.04)	0.485($\pm$0.02)	0.072($\pm$0.04)
	Gerry Fair Classifier	0.371($\pm$0.04)	0.423($\pm$0.02)	0.052($\pm$0.03)
COMPAS Recidivism	Fair Transition Loss	0.208($\pm$0.06)	0.283($\pm$0.02)	0.074($\pm$0.05)
	Adversarial Debiasing	0.191($\pm$0.11)	0.324($\pm$0.03)	0.133($\pm$0.10)
	Gerry Fair Classifier	0.155($\pm$0.05)	0.274($\pm$0.06)	0.120($\pm$0.04)
	Adaptive Priority Reweighting	$-$0.111($\pm$0.18)	0.260($\pm$0.04)	0.371($\pm$0.17)
	Prejudice Remover	$-$0.352($\pm$0.03)	$-$0.278($\pm$0.02)	0.073($\pm$0.03)
	Standard MLP (baseline)	$-$0.471($\pm$0.05)	$-$0.294($\pm$0.02)	0.176($\pm$0.04)
German Credit	Adaptive Priority Reweighting	0.299($\pm$0.08)	0.373($\pm$0.06)	0.075($\pm$0.05)
	Prejudice Remover	0.283($\pm$0.10)	0.391($\pm$0.07)	0.107($\pm$0.06)
	Fair Transition Loss	0.273($\pm$0.10)	0.386($\pm$0.08)	0.113($\pm$0.08)
	Standard MLP (baseline)	0.270($\pm$0.07)	0.352($\pm$0.05)	0.082($\pm$0.04)
	Gerry Fair Classifier	0.218($\pm$0.11)	0.321($\pm$0.10)	0.103($\pm$0.05)
	Adversarial Debiasing	0.040($\pm$0.41)	0.301($\pm$0.13)	0.261($\pm$0.30)

Tabela A.3: Complete results optimizing MCC and Equalized Odds.

Dataset	Method	Fitness	MCC	Eq. Odds
Adult Income	Fair Transition Loss	0.556($\pm$0.03)	0.584($\pm$0.01)	0.029($\pm$0.03)
	Adaptive Priority Reweighting	0.553($\pm$0.01)	0.576($\pm$0.01)	0.022($\pm$0.02)
	Prejudice Remover	0.505($\pm$0.09)	0.560($\pm$0.02)	0.055($\pm$0.08)
	Adversarial Debiasing	0.493($\pm$0.05)	0.573($\pm$0.01)	0.080($\pm$0.05)
	Standard MLP (baseline)	0.489($\pm$0.03)	0.580($\pm$0.01)	0.091($\pm$0.03)
	Gerry Fair Classifier	0.394($\pm$0.03)	0.515($\pm$0.02)	0.121($\pm$0.02)
Bank Marketing	Fair Transition Loss	0.467($\pm$0.11)	0.560($\pm$0.03)	0.093($\pm$0.10)
	Adaptive Priority Reweighting	0.441($\pm$0.06)	0.500($\pm$0.01)	0.059($\pm$0.06)
	Standard MLP (baseline)	0.432($\pm$0.06)	0.520($\pm$0.02)	0.087($\pm$0.06)
	Prejudice Remover	0.392($\pm$0.09)	0.490($\pm$0.02)	0.098($\pm$0.08)
	Adversarial Debiasing	0.373($\pm$0.09)	0.508($\pm$0.02)	0.136($\pm$0.09)
	Gerry Fair Classifier	0.344($\pm$0.07)	0.422($\pm$0.02)	0.078($\pm$0.06)
COMPAS Recidivism	Adaptive Priority Reweighting	0.292($\pm$0.03)	0.319($\pm$0.02)	0.027($\pm$0.02)
	Adversarial Debiasing	0.258($\pm$0.05)	0.329($\pm$0.03)	0.070($\pm$0.05)
	Fair Transition Loss	0.213($\pm$0.06)	0.264($\pm$0.06)	0.050($\pm$0.03)
	Gerry Fair Classifier	0.201($\pm$0.05)	0.290($\pm$0.04)	0.089($\pm$0.05)
	Prejudice Remover	$-0.319(\pm0.03)$	$-0.289(\pm0.03)$	0.030($\pm$0.02)
	Standard MLP (baseline)	$-0.435(\pm0.03)$	$-0.292(\pm0.02)$	0.143($\pm$0.03)
German Credit	Standard MLP (baseline)	0.295($\pm$0.09)	0.354($\pm$0.08)	0.060($\pm$0.04)
	Fair Transition Loss	0.274($\pm$0.10)	0.361($\pm$0.08)	0.087($\pm$0.05)
	Gerry Fair Classifier	0.273($\pm$0.10)	0.361($\pm$0.06)	0.087($\pm$0.06)
	Prejudice Remover	0.271($\pm$0.07)	0.324($\pm$0.06)	0.054($\pm$0.04)
	Adaptive Priority Reweighting	0.261($\pm$0.08)	0.326($\pm$0.06)	0.065($\pm$0.05)
	Adversarial Debiasing	0.116($\pm$0.40)	0.311($\pm$0.14)	0.195($\pm$0.28)

Tabela A.4: Complete results optimizing Accuracy and Statistical Parity.

Dataset	Method	Fitness	Accuracy	Stat. Parity
Adult Income	Fair Transition Loss	0.814($\pm$0.01)	0.828($\pm$0.01)	0.014($\pm$0.01)
	Adaptive Priority Reweighting	0.811($\pm$0.01)	0.822($\pm$0.01)	0.011($\pm$0.01)
	Adversarial Debiasing	0.808($\pm$0.01)	0.830($\pm$0.01)	0.022($\pm$0.01)
	Prejudice Remover	0.807($\pm$0.01)	0.825($\pm$0.00)	0.018($\pm$0.01)
	Standard MLP (baseline)	0.666($\pm$0.01)	0.851($\pm$0.00)	0.184($\pm$0.01)
	Gerry Fair Classifier	0.651($\pm$0.21)	0.721($\pm$0.07)	0.070($\pm$0.14)
Bank Marketing	Adversarial Debiasing	0.869($\pm$0.03)	0.901($\pm$0.00)	0.031($\pm$0.02)
	Prejudice Remover	0.860($\pm$0.02)	0.898($\pm$0.00)	0.038($\pm$0.02)
	Fair Transition Loss	0.854($\pm$0.05)	0.889($\pm$0.01)	0.035($\pm$0.05)
	Adaptive Priority Reweighting	0.851($\pm$0.06)	0.900($\pm$0.00)	0.049($\pm$0.06)
	Gerry Fair Classifier	0.824($\pm$0.02)	0.895($\pm$0.00)	0.071($\pm$0.02)
	Standard MLP (baseline)	0.799($\pm$0.04)	0.902($\pm$0.00)	0.103($\pm$0.03)
COMPAS Recidivism	Adversarial Debiasing	0.538($\pm$0.07)	0.670($\pm$0.02)	0.132($\pm$0.08)
	Fair Transition Loss	0.501($\pm$0.15)	0.600($\pm$0.05)	0.099($\pm$0.14)
	Gerry Fair Classifier	0.501($\pm$0.05)	0.614($\pm$0.05)	0.113($\pm$0.07)
	Prejudice Remover	0.308($\pm$0.03)	0.359($\pm$0.01)	0.052($\pm$0.02)
	Standard MLP (baseline)	0.146($\pm$0.03)	0.354($\pm$0.02)	0.208($\pm$0.02)
	Adaptive Priority Reweighting	$-$0.105($\pm$0.26)	0.584($\pm$0.03)	0.689($\pm$0.23)
German Credit	Prejudice Remover	0.684($\pm$0.05)	0.757($\pm$0.02)	0.073($\pm$0.06)
	Standard MLP (baseline)	0.639($\pm$0.06)	0.752($\pm$0.02)	0.113($\pm$0.06)
	Gerry Fair Classifier	0.621($\pm$0.09)	0.712($\pm$0.12)	0.090($\pm$0.04)
	Fair Transition Loss	0.616($\pm$0.20)	0.715($\pm$0.06)	0.098($\pm$0.17)
	Adaptive Priority Reweighting	0.589($\pm$0.08)	0.682($\pm$0.03)	0.093($\pm$0.08)
	Adversarial Debiasing	0.430($\pm$0.33)	0.713($\pm$0.09)	0.283($\pm$0.26)

Tabela A.5: Complete results optimizing Accuracy and Equal Opportunity.

Dataset	Method	Fitness	Accuracy	Eq. Opp.
Adult Income	Adaptive Priority Reweighting	0.808($\pm$0.01)	0.837($\pm$0.00)	0.029($\pm$0.01)
	Fair Transition Loss	0.808($\pm$0.02)	0.842($\pm$0.01)	0.034($\pm$0.02)
	Adversarial Debiasing	0.796($\pm$0.01)	0.849($\pm$0.00)	0.052($\pm$0.01)
	Prejudice Remover	0.794($\pm$0.02)	0.845($\pm$0.01)	0.051($\pm$0.01)
	Standard MLP (baseline)	0.765($\pm$0.02)	0.850($\pm$0.00)	0.084($\pm$0.02)
	Gerry Fair Classifier	0.756($\pm$0.01)	0.788($\pm$0.03)	0.032($\pm$0.04)
Bank Marketing	Adaptive Priority Reweighting	0.858($\pm$0.02)	0.897($\pm$0.00)	0.039($\pm$0.03)
	Gerry Fair Classifier	0.837($\pm$0.04)	0.895($\pm$0.00)	0.058($\pm$0.04)
	Fair Transition Loss	0.833($\pm$0.05)	0.892($\pm$0.01)	0.059($\pm$0.05)
	Prejudice Remover	0.827($\pm$0.04)	0.898($\pm$0.00)	0.071($\pm$0.04)
	Standard MLP (baseline)	0.826($\pm$0.04)	0.901($\pm$0.00)	0.075($\pm$0.04)
	Adversarial Debiasing	0.807($\pm$0.07)	0.902($\pm$0.00)	0.095($\pm$0.07)
COMPAS Recidivism	Fair Transition Loss	0.572($\pm$0.03)	0.631($\pm$0.04)	0.059($\pm$0.03)
	Adversarial Debiasing	0.553($\pm$0.09)	0.669($\pm$0.01)	0.116($\pm$0.09)
	Gerry Fair Classifier	0.530($\pm$0.03)	0.637($\pm$0.04)	0.107($\pm$0.05)
	Adaptive Priority Reweighting	0.356($\pm$0.18)	0.643($\pm$0.02)	0.287($\pm$0.18)
	Prejudice Remover	0.264($\pm$0.03)	0.357($\pm$0.01)	0.093($\pm$0.02)
	Standard MLP (baseline)	0.155($\pm$0.04)	0.350($\pm$0.02)	0.195($\pm$0.04)
German Credit	Adaptive Priority Reweighting	0.674($\pm$0.06)	0.750($\pm$0.03)	0.076($\pm$0.04)
	Prejudice Remover	0.664($\pm$0.05)	0.748($\pm$0.02)	0.084($\pm$0.04)
	Gerry Fair Classifier	0.662($\pm$0.12)	0.719($\pm$0.12)	0.057($\pm$0.07)
	Standard MLP (baseline)	0.638($\pm$0.06)	0.738($\pm$0.04)	0.101($\pm$0.05)
	Fair Transition Loss	0.599($\pm$0.12)	0.711($\pm$0.05)	0.112($\pm$0.11)
	Adversarial Debiasing	0.368($\pm$0.38)	0.685($\pm$0.10)	0.317($\pm$0.30)

Tabela A.6: Complete results optimizing Accuracy and Equalized Odds.

Dataset	Method	Fitness	Accuracy	Eq. Odds
Adult Income	Adaptive Priority Reweighting	0.829(±0.01)	0.847(±0.00)	0.018(±0.01)
	Prejudice Remover	0.810(±0.02)	0.846(±0.00)	0.036(±0.02)
	Fair Transition Loss	0.787(±0.08)	0.826(±0.07)	0.039(±0.04)
	Adversarial Debiasing	0.756(±0.03)	0.848(±0.00)	0.092(±0.03)
	Standard MLP (baseline)	0.752(±0.04)	0.849(±0.00)	0.097(±0.04)
	Gerry Fair Classifier	0.705(±0.07)	0.751(±0.09)	0.046(±0.05)
Bank Marketing	Adaptive Priority Reweighting	0.846(±0.05)	0.901(±0.00)	0.055(±0.05)
	Gerry Fair Classifier	0.837(±0.06)	0.893(±0.00)	0.057(±0.06)
	Standard MLP (baseline)	0.800(±0.06)	0.902(±0.00)	0.102(±0.06)
	Fair Transition Loss	0.799(±0.10)	0.891(±0.01)	0.092(±0.10)
	Prejudice Remover	0.781(±0.07)	0.899(±0.00)	0.118(±0.07)
	Adversarial Debiasing	0.750(±0.09)	0.900(±0.00)	0.150(±0.09)
COMPAS Recidivism	Adaptive Priority Reweighting	0.642(±0.03)	0.669(±0.01)	0.027(±0.02)
	Fair Transition Loss	0.594(±0.04)	0.648(±0.01)	0.054(±0.03)
	Adversarial Debiasing	0.594(±0.07)	0.672(±0.02)	0.078(±0.06)
	Gerry Fair Classifier	0.558(±0.05)	0.647(±0.02)	0.088(±0.04)
	Prejudice Remover	0.287(±0.03)	0.342(±0.01)	0.055(±0.03)
	Standard MLP (baseline)	0.218(±0.05)	0.353(±0.01)	0.135(±0.05)
German Credit	Adaptive Priority Reweighting	0.716(±0.04)	0.750(±0.02)	0.034(±0.03)
	Standard MLP (baseline)	0.681(±0.08)	0.747(±0.03)	0.066(±0.06)
	Prejudice Remover	0.648(±0.06)	0.743(±0.03)	0.095(±0.06)
	Gerry Fair Classifier	0.643(±0.13)	0.707(±0.13)	0.063(±0.04)
	Fair Transition Loss	0.622(±0.26)	0.705(±0.10)	0.083(±0.17)
	Adversarial Debiasing	0.530(±0.33)	0.713(±0.10)	0.183(±0.24)

Capítulo 5

Conclusions

In this study, we present Fair Transition Loss, a novel in-processing technique for addressing fair classification problems. It leverages concepts from label noise robustness to mitigate social bias against underprivileged groups. We delve into the intersection of these two research areas, highlighting both their similarities and differences. Our approach tackles the fairness-performance trade-off as a multi-objective optimization problem, employing a linear relaxed objective function to reduce bias while maintaining acceptable predictive performance levels. We benchmark this approach and compare to proeminent in-processing techniques in common fair classification tasks, using the Almost Stochastic Order test to evaluate results through multiple resampling iterations. This ensures that all methods operate under the same conditions, maximizing their potential within the scope of hyperparameter tuning.

This is the first technique that models fair classification problems by drawing insights from classification in the presence of label noise. Our experiments indicate that Fair Transition Loss consistently outperforms its competitors in most optimization scenarios. Even in those cases that the proposed method isn't the outright leader, it performs at least as well as evaluated alternatives. Therefore, this novel approach can significantly mitigate bias while keeping model performance, particularly in scenarios optimizing balanced performance metrics like MCC. The proposed technique particularly stands out in setups where hyperparameter tuning is an integral component of the prediction pipeline.

While our proposed method seems competitive in problems involving hyperparameter optimization for binary fair classification tasks using a simple Multi-Layer Perceptron, we can outline some potential research directions: evaluate Fair Transition Loss within different neural network architectures, such as Deep Neural Networks; investigate whether the proposed method can effectively address multi-class fair classification problems and handle multiple sensitive attributes, as theoretically possible; evaluate FTL within different multi objective optimization schemes, such as the Fair

Hyperparameter Tuning techniques proposed by [44] or the non-linear Chebyshev scalarization scheme proposed by [21]; explore approaches to estimating or initializing transition matrices without relying on hyperparameter tuning techniques.

With this work, we hope to establish Fair Transition Loss as a valuable tool in fair classification tasks and pave the way for novel approaches that draw insights from label noise for various fair machine learning problems, including regression, recommender systems, ranking and language models.

Referências Bibliográficas

[1] MEHRABI, N., MORSTATTER, F., SAXENA, N., et al. "A Survey on Bias and Fairness in Machine Learning", *ACM Comput. Surv.*, v. 54, n. 6, jul 2021. ISSN: 0360-0300. doi: 10.1145/3457607. Disponível em: <https://doi.org/10.1145/3457607>.

[2] HUTCHINSON, B., MITCHELL, M. "50 Years of Test (Un)fairness: Lessons for Machine Learning", *Proceedings of the Conference on Fairness, Accountability, and Transparency*, 2018. Disponível em: <https://api.semanticscholar.org/CorpusID:53782832>.

[3] CATON, S., HAAS, C. "Fairness in Machine Learning: A Survey", *ACM Comput. Surv.*, aug 2023. ISSN: 0360-0300. doi: 10.1145/3616865. Disponível em: <https://doi.org/10.1145/3616865>. Just Accepted.

[4] ZAFAR, M. B., VALERA, I., ROGRIGUEZ, M. G., et al. "Fairness Constraints: Mechanisms for Fair Classification". In: Singh, A., Zhu, J. (Eds.), *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics*, v. 54, *Proceedings of Machine Learning Research*, pp. 962–970. PMLR, 20–22 Apr 2017. Disponível em: <https://proceedings.mlr.press/v54/zafar17a.html>.

[5] VERMA, S., RUBIN, J. "Fairness Definitions Explained", *IEEE/ACM International Workshop on Software Fairness*, v. 18, 2018. doi: 10.1145/3194770.3194776. Disponível em: <https://doi.org/10.1145/3194770.3194776>.

[6] ALER TUBELLA, A., BARSOTTI, F., KOÇER, R. G., et al. "Ethical implications of fairness interventions: what might be hidden behind engineering choices?" *Ethics and Information Technology*, v. 24, n. 1, pp. 1–11, mar 2022. ISSN: 15728439. doi: 10.1007/S10676-022-09636-Z/TABLES/4. Disponível em: <https://link.springer.com/article/10.1007/s10676-022-09636-z>.

[7] ALVES, G., BERNIER, F., COUCEIRO, M., et al. "Survey on fairness notions and related tensions", *EURO Journal on Decision Processes*, v. 11, pp. 100033, 2023. ISSN: 2193-9438. doi: https://doi.org/10.1016/j.ejdp.2023.100033. Disponível em: <`https://www.sciencedirect.com/science/article/pii/S2193943823000067`>.

[8] WEINBERG, L. "Rethinking Fairness: An Interdisciplinary Survey of Critiques of Hegemonic ML Fairness Approaches", *Journal of Artificial Intelligence Research*, v. 74, pp. 75–109, 2022. doi: 10.1613/jair.1.13196.

[9] DWORK, C., HARDT, M., PITASSI, T., et al. "Fairness through Awareness". In: *Proceedings of the 3rd Innovations in Theoretical Computer Science Conference*, ITCS '12, p. 214–226, New York, NY, USA, 2012. Association for Computing Machinery. ISBN: 9781450311151. doi: 10.1145/2090236.2090255. Disponível em: <`https://doi.org/10.1145/2090236.2090255`>.

[10] KUSNER, M., LOFTUS, J., RUSSELL, C., et al. "Counterfactual Fairness". In: *Proceedings of the 31st International Conference on Neural Information Processing Systems*, NIPS'17, p. 4069–4079, Red Hook, NY, USA, 2017. Curran Associates Inc. ISBN: 9781510860964.

[11] HARDT, M., PRICE, E., SREBRO, N. "Equality of Opportunity in Supervised Learning". In: *Proceedings of the 30th International Conference on Neural Information Processing Systems*, NIPS'16, p. 3323–3331, Red Hook, NY, USA, 2016. Curran Associates Inc. ISBN: 9781510838819.

[12] CALDERS, T., VERWER, S. "Three naive Bayes approaches for discrimination-free classification", *Data Mining and Knowledge Discovery*, v. 21, pp. 277–292, 2010. ISSN: 13845810. doi: 10.1007/s10618-010-0190-x.

[13] KAMISHIMA, T., AKAHO, S., ASOH, H., et al. "Fairness-aware classifier with prejudice remover regularizer", *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*, v. 7524 LNAI, n. PART 2, pp. 35–50, 2012. ISSN: 03029743. doi: 10.1007/978-3-642-33486-3.

[14] WOODWORTH, B., GUNASEKAR, S., OHANNESSIAN, M. I., et al. "Learning Non-Discriminatory Predictors". In: Kale, S., Shamir, O. (Eds.), *Proceedings of the 2017 Conference on Learning Theory*, v. 65, *Proceedings of Machine Learning Research*, pp. 1920–1953. PMLR, 07–10

Jul 2017. Disponível em: <https://proceedings.mlr.press/v65/woodworth17a.html>.

[15] ZAFAR, M. B., VALERA, I., GOMEZ RODRIGUEZ, M., et al. "Fairness Beyond Disparate Treatment & Disparate Impact: Learning Classification without Disparate Mistreatment". In: *Proceedings of the 26th International Conference on World Wide Web*, WWW '17, p. 1171–1180, Republic and Canton of Geneva, CHE, 2017. International World Wide Web Conferences Steering Committee. ISBN: 9781450349130. doi: 10.1145/3038912.3052660. Disponível em: <https://doi.org/10.1145/3038912.3052660>.

[16] KEARNS, M., NEEL, S., ROTH, A., et al. "Preventing Fairness Gerrymandering: Auditing and Learning for Subgroup Fairness". In: Dy, J., Krause, A. (Eds.), *Proceedings of the 35th International Conference on Machine Learning*, v. 80, *Proceedings of Machine Learning Research*, pp. 2564–2572. PMLR, 10–15 Jul 2018. Disponível em: <https://proceedings.mlr.press/v80/kearns18a.html>.

[17] KEARNS, M., NEEL, S., ROTH, A., et al. "An Empirical Study of Rich Subgroup Fairness for Machine Learning". In: *Proceedings of the Conference on Fairness, Accountability, and Transparency*, FAT* '19, p. 100–109, New York, NY, USA, 2019. Association for Computing Machinery. ISBN: 9781450361255. doi: 10.1145/3287560.3287592. Disponível em: <https://doi.org/10.1145/3287560.3287592>.

[18] ADEL, T., VALERA, I., GHAHRAMANI, Z., et al. "One-Network Adversarial Fairness", *Proceedings of the AAAI Conference on Artificial Intelligence*, v. 33, n. 01, pp. 2412–2420, Jul. 2019. doi: 10.1609/aaai.v33i01.33012412. Disponível em: <https://ojs.aaai.org/index.php/AAAI/article/view/4085>.

[19] XU, D., YUAN, S., ZHANG, L., et al. "FairGAN+: Achieving Fair Data Generation and Classification through Generative Adversarial Nets", *Proceedings - 2019 IEEE International Conference on Big Data, Big Data 2019*, pp. 1401–1406, 2019. doi: 10.1109/BigData47090.2019.9006322.

[20] PETROVIĆ, A., NIKOLIĆ, M., JOVANOVIĆ, M., et al. "Fair classification via Monte Carlo policy gradient method", *Engineering Applications of Artificial Intelligence*, v. 104, n. February, pp. 104398, 2021. ISSN: 09521976. doi: 10.1016/j.engappai.2021.104398. Disponível em: <https://doi.org/10.1016/j.engappai.2021.104398>.

[21] WEI, S., NIETHAMMER, M. "The fairness-accuracy Pareto front", *Statistical Analysis and Data Mining*, v. 15, pp. 287–302, 6 2022. ISSN: 19321872. doi: 10.1002/SAM.11560.

[22] LIU, S., VICENTE, L. N. "Accuracy and fairness trade-offs in machine learning: a stochastic multi-objective approach", *Computational Management Science*, v. 19, pp. 513–537, 7 2022. ISSN: 16196988. doi: 10.1007/s10287-022-00425-z.

[23] MERCIER, Q., POIRION, F., DÉSIDÉRI, J. A. "A stochastic multiple gradient descent algorithm", *European Journal of Operational Research*, v. 271, pp. 808–817, 12 2018. ISSN: 03772217. doi: 10.1016/j.ejor.2018.05.064.

[24] HU, Z., XU, Y., TIAN, X. "Adaptive Priority Reweighing for Generalizing Fairness Improvement". In: *International Joint Conference on Neural Networks, IJCNN 2023, Gold Coast, Australia, June 18-23, 2023*, pp. 1–8. IEEE, 2023. doi: 10.1109/IJCNN54540.2023.10191757. Disponível em: <https://doi.org/10.1109/IJCNN54540.2023.10191757>.

[25] D'ALOISIO, G., D'ANGELO, A., DI MARCO, A., et al. "Debiaser for Multiple Variables to enhance fairness in classification tasks", *Information Processing and Management*, v. 60, n. 2, pp. 103226, 2023. ISSN: 03064573. doi: 10.1016/j.ipm.2022.103226. Disponível em: <https://doi.org/10.1016/j.ipm.2022.103226>.

[26] LIU, T., WANG, H., WANG, Y., et al. "SimFair: A Unified Framework for Fairness-Aware Multi-Label Classification", *Proceedings of the AAAI Conference on Artificial Intelligence*, v. 37, n. 12, pp. 14338–14346, 2023. doi: 10.1609/aaai.v37i12.26677. Disponível em: <https://ojs.aaai.org/index.php/AAAI/article/view/26677>.

[27] KIM, D., PARK, S., HWANG, S., et al. "Fair classification by loss balancing via fairness-aware batch sampling", *Neurocomputing*, v. 518, pp. 231–241, 2023. ISSN: 0925-2312. doi: https://doi.org/10.1016/j.neucom.2022.11.018. Disponível em: <https://www.sciencedirect.com/science/article/pii/S0925231222013984>.

[28] KHALILI, M. M., ZHANG, X., ABROSHAN, M. "Loss Balancing for Fair Supervised Learning". In: Krause, A., Brunskill, E., Cho, K., et al. (Eds.), *International Conference on Machine Learning, ICML 2023, 23-29 July 2023, Honolulu, Hawaii, USA*, v. 202, *Proceedings of Machine Learning Research*, pp. 16271–16290. PMLR, 2023. Disponível em: <https://proceedings.mlr.press/v202/khalili23a.html>.

[29] LIANG, Y., CHEN, C., TIAN, T., et al. "Fair classification via domain adaptation: A dual adversarial learning approach", *Frontiers in Big Data*, v. 5, 2023. ISSN: 2624-909X. doi: 10.3389/fdata.2022.1049565. Disponível em: <https://www.frontiersin.org/articles/10.3389/fdata.2022.1049565>.

[30] ZHANG, T., ZHU, T., LI, J., et al. "Revisiting model fairness via adversarial examples", *Knowl. Based Syst.*, v. 277, pp. 110777, 2023. doi: 10.1016/J.KNOSYS.2023.110777. Disponível em: <https://doi.org/10.1016/j.knosys.2023.110777>.

[31] GRARI, V., LAMPRIER, S., DETYNIECKI, M. "Adversarial learning for counterfactual fairness", *Mach. Learn.*, v. 112, n. 3, pp. 741–763, 2023. doi: 10.1007/S10994-022-06206-8. Disponível em: <https://doi.org/10.1007/s10994-022-06206-8>.

[32] MOUSAVI, S. A., MOUSAVI, H., DANESHTALAB, M. "FARMUR: Fair Adversarial Retraining to Mitigate Unfairness in Robustness". In: Abelló, A., Vassiliadis, P., Romero, O., et al. (Eds.), *Advances in Databases and Information Systems - 27th European Conference, ADBIS 2023, Barcelona, Spain, September 4-7, 2023, Proceedings*, v. 13985, *Lecture Notes in Computer Science*, pp. 133–145. Springer, 2023. doi: 10.1007/978-3-031-42914-9_10. Disponível em: <https://doi.org/10.1007/978-3-031-42914-9_10>.

[33] WEI, Z., WANG, Y., GUO, Y., et al. "CFA: Class-wise Calibrated Fair Adversarial Training", *CoRR*, v. abs/2303.14460, 2023. doi: 10.48550/ARXIV.2303.14460. Disponível em: <https://doi.org/10.48550/arXiv.2303.14460>.

[34] MA, Y., FRAUEN, D., MELNYCHUK, V., et al. "Counterfactual Fairness for Predictions using Generative Adversarial Networks", *CoRR*, v. abs/2310.17687, 2023. doi: 10.48550/ARXIV.2310.17687. Disponível em: <https://doi.org/10.48550/arXiv.2310.17687>.

[35] CHEN, H., ZHU, T., ZHANG, T., et al. "Privacy and Fairness in Federated Learning: On the Perspective of Tradeoff", *ACM Comput. Surv.*, v. 56, n. 2, pp. 39:1–39:37, 2024. doi: 10.1145/3606017. Disponível em: <https://doi.org/10.1145/3606017>.

[36] VUCINICH, S., ZHU, Q. "The Current State and Challenges of Fairness in Federated Learning", *IEEE Access*, v. 11, pp. 80903–80914, 2023. doi:

10.1109/ACCESS.2023.3295412. Disponível em: <https://doi.org/10.1109/ACCESS.2023.3295412>.

[37] ZHANG, W., WEISS, J. C. "Longitudinal Fairness with Censorship". In: *Thirty-Sixth AAAI Conference on Artificial Intelligence, AAAI 2022, Thirty-Fourth Conference on Innovative Applications of Artificial Intelligence, IAAI 2022, The Twelveth Symposium on Educational Advances in Artificial Intelligence, EAAI 2022 Virtual Event, February 22 - March 1, 2022*, pp. 12235–12243. AAAI Press, 2022. doi: 10.1609/AAAI.V36I11.21484. Disponível em: <https://doi.org/10.1609/aaai.v36i11.21484>.

[38] ZHANG, W., HERNANDEZ-BOUSSARD, T., WEISS, J. "Censored Fairness through Awareness", *Proceedings of the AAAI Conference on Artificial Intelligence*, v. 37, n. 12, pp. 14611–14619, Jun. 2023. doi: 10.1609/aaai.v37i12.26708. Disponível em: <https://ojs.aaai.org/index.php/AAAI/article/view/26708>.

[39] ZHANG, W., WEISS, J. C. "Fairness with censorship and group constraints", *Knowl. Inf. Syst.*, v. 65, n. 6, pp. 2571–2594, 2023. doi: 10.1007/S10115-023-01842-5. Disponível em: <https://doi.org/10.1007/s10115-023-01842-5>.

[40] ZHANG, W., KIM, J., WANG, Z., et al. "Individual Fairness Guarantee in Learning with Censorship", *CoRR*, v. abs/2302.08015, 2023. doi: 10.48550/ARXIV.2302.08015. Disponível em: <https://doi.org/10.48550/arXiv.2302.08015>.

[41] PARETO, V. "Manuale di economica politica, societa editrice libraria", *Manual of political economy*, v. 1971, 1906.

[42] SCHMUCKER, R., DONINI, M., PERRONE, V., et al. "Multi-objective multi-fidelity hyperparameter optimization with application to fairness". In: *NeurIPS 2020 Workshop on Meta-learning*, 2020.

[43] LI, L., JAMIESON, K., DESALVO, G., et al. "Hyperband: A Novel Bandit-Based Approach to Hyperparameter Optimization", *J. Mach. Learn. Res.*, v. 18, n. 1, pp. 6765–6816, jan 2017. ISSN: 1532-4435.

[44] F.CRUZ, A., SALEIRO, P., BELÉM, C., et al. "Promoting Fairness through Hyperparameter Optimization". In: *2021 IEEE International Conference on Data Mining (ICDM)*, pp. 1036–1041, 2021. doi: 10.1109/ICDM51629.2021.00119.

[45] GIAGKIOZIS, I., FLEMING, P. J. "Methods for multi-objective optimization: An analysis", *Information Sciences*, v. 293, pp. 338–350, 2015. ISSN: 0020-0255. doi: https://doi.org/10.1016/j.ins.2014.08.071. Disponível em: `<https://www.sciencedirect.com/science/article/pii/S0020025514009074>`.

[46] GOH, G., COTTER, A., GUPTA, M., et al. "Satisfying Real-world Goals with Dataset Constraints". In: Lee, D., Sugiyama, M., Luxburg, U., et al. (Eds.), *Advances in Neural Information Processing Systems*, v. 29. Curran Associates, Inc., 2016. Disponível em: `<https://proceedings.neurips.cc/paper_files/paper/2016/file/dc4c44f624d600aa568390f1f1104aa0-Paper.pdf>`.

[47] FRÉNAY, B., VERLEYSEN, M. "Classification in the presence of label noise: A survey", *IEEE Transactions on Neural Networks and Learning Systems*, v. 25, n. 5, pp. 845–869, 2014. ISSN: 21622388. doi: 10.1109/TNNLS.2013.2292894.

[48] HICKEY, R. J. "Noise modelling and evaluating learning from examples", *Artificial Intelligence*, v. 82, n. 1–2, pp. 157–179, apr 1996. ISSN: 0004-3702. doi: http://dx.doi.org/10.1016/0004-3702(94)00094-8. Disponível em: `<http://www.sciencedirect.com/science/article/pii/0004370294000948>`.

[49] QUINLAN, J. R. "Induction of decision trees", *Machine Learning 1986 1:1*, v. 1, n. 1, pp. 81–106, mar 1986. ISSN: 1573-0565. doi: 10.1007/BF00116251. Disponível em: `<https://link.springer.com/article/10.1007/BF00116251>`.

[50] PATRINI, G., ROZZA, A., MENON, A. K., et al. "Making deep neural networks robust to label noise: A loss correction approach", *Proceedings - 30th IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2017*, v. 2017-Janua, pp. 2233–2241, 2017. doi: 10.1109/CVPR.2017.240.

[51] PATRINI, G., NIELSEN, F., NOCK, R., et al. "Loss Factorization, Weakly Supervised Learning and Label Noise Robustness". In: *Proceedings of the 33rd International Conference on International Conference on Machine Learning - Volume 48*, ICML'16, p. 708–717. JMLR.org, 2016.

[52] LAMY, A., ZHONG, Z., VERMA, N., et al. "Noise-tolerant fair classification", *Advances in Neural Information Processing Systems*, v. 32, 1 2019. ISSN: 10495258. doi: 10.48550/arxiv.1901.10837. Disponível em: `<https://arxiv.org/abs/1901.10837v4>`.

[53] FOGLIATO, R., CHOULDECHOVA, A., G'SELL, M. "Fairness Evaluation in Presence of Biased Noisy Labels". In: Chiappa, S., Calandra, R. (Eds.), *Proceedings of the Twenty Third International Conference on Artificial Intelligence and Statistics*, v. 108, *Proceedings of Machine Learning Research*, pp. 2325–2336. PMLR, 26–28 Aug 2020. Disponível em: <https://proceedings.mlr.press/v108/fogliato20a.html>.

[54] WANG, S., GUO, W., NARASIMHAN, H., et al. "Robust Optimization for Fairness with Noisy Protected Groups", *Advances in Neural Information Processing Systems*, v. 2020-December, feb 2020. ISSN: 10495258. doi: 10.48550/arxiv.2002.09343. Disponível em: <https://arxiv.org/abs/2002.09343v3>.

[55] MEHROTRA, A., CELIS, L. E. "Mitigating Bias in Set Selection with Noisy Protected Attributes", *FAccT 2021 - Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, pp. 237–248, nov 2021. doi: 10.48550/arxiv.2011.04219. Disponível em: <https://arxiv.org/abs/2011.04219v2>.

[56] CELIS, L. E., HUANG, L., KESWANI, V., et al. "Fair Classification with Noisy Protected Attributes: A Framework with Provable Guarantees". In: Meila, M., Zhang, T. (Eds.), *Proceedings of the 38th International Conference on Machine Learning*, v. 139, *Proceedings of Machine Learning Research*, pp. 1349–1361. PMLR, 18–24 Jul 2021. Disponível em: <https://proceedings.mlr.press/v139/celis21a.html>.

[57] PROST, F., AWASTHI, P., BLUMM, N., et al. "Measuring Model Fairness under Noisy Covariates: A Theoretical Perspective". In: *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society*, AIES '21, p. 873–883, New York, NY, USA, 2021. Association for Computing Machinery. ISBN: 9781450384735. doi: 10.1145/3461702.3462603. Disponível em: <https://doi.org/10.1145/3461702.3462603>.

[58] GHAZIMATIN, A., KLEINDESSNER, M., RUSSELL, C., et al. "Measuring Fairness of Rankings under Noisy Sensitive Information". In: *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency*, FAccT '22, p. 2263–2279, New York, NY, USA, 2022. Association for Computing Machinery. ISBN: 9781450393522. doi: 10.1145/3531146.3534641. Disponível em: <https://doi.org/10.1145/3531146.3534641>.

[59] ZHANG, T., ZHU, T., LI, J., et al. "Fairness in Semi-Supervised Learning: Unlabeled Data Help to Reduce Discrimination", *IEEE Transactions on Knowledge and Data Engineering*, v. 34, n. 4, pp. 1763–1774, apr 2022. ISSN: 15582191. doi: 10.1109/TKDE.2020.3002567.

[60] WANG, J., CRUZ, U. S. C. S., LIU, U. Y., et al. "Fair classification with group-dependent label noise", *FAccT 2021 - Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, pp. 526–536, 3 2021. doi: 10.1145/3442188.3445915. Disponível em: <https://doi.org/10.1145/3442188.3445915>.

[61] WU, S., HAN, B., LIU, Y., et al. "Fair Classification with Instance-dependent Label Noise", *Proceedings of Machine Learning Research*, v. 140, pp. 1–17, 2022. Disponível em: <https://www.mturk.com/>.

[62] GHOSH, A., KVITCA, P., WILSON, C. "When Fair Classification Meets Noisy Protected Attributes". In: *Proceedings of the 2023 AAAI/ACM Conference on AI, Ethics, and Society*, AIES '23, p. 679–690, New York, NY, USA, 2023. Association for Computing Machinery. ISBN: 9798400702310. doi: 10.1145/3600211.3604707. Disponível em: <https://doi.org/10.1145/3600211.3604707>.

[63] ZHANG, B. H., LEMOINE, B., MITCHELL, M. "Mitigating Unwanted Biases with Adversarial Learning". pp. 335–340. Association for Computing Machinery, Inc, 12 2018. ISBN: 9781450360128. doi: 10.1145/3278721.3278779.

[64] BELLAMY, R. K. E., DEY, K., HIND, M., et al. "AI Fairness 360: An Extensible Toolkit for Detecting, Understanding, and Mitigating Unwanted Algorithmic Bias". out. 2018. Disponível em: <https://www.ibm.com/opensource/open/projects/ai-fairness-360/>.

[65] LI, M., SOLTANOLKOTABI, M., OYMAK, S. "Gradient Descent with Early Stopping is Provably Robust to Label Noise for Overparameterized Neural Networks". In: Chiappa, S., Calandra, R. (Eds.), *Proceedings of the Twenty Third International Conference on Artificial Intelligence and Statistics*, v. 108, *Proceedings of Machine Learning Research*, pp. 4313–4324. PMLR, 26–28 Aug 2020. Disponível em: <https://proceedings.mlr.press/v108/li20j.html>.

[66] KINGMA, D. P., BA, J. "Adam: A Method for Stochastic Optimization". In: Bengio, Y., LeCun, Y. (Eds.), *3rd International Conference on*

Learning Representations, ICLR 2015, San Diego, CA, USA, May 7-9, 2015, Conference Track Proceedings, 2015. Disponível em: <`http://arxiv.org/abs/1412.6980`>.

[67] JIANG, H., NACHUM, O. "Identifying and correcting label bias in machine learning". In: *International Conference on Artificial Intelligence and Statistics*, pp. 702–712. PMLR, 2020.

[68] MROUEH, Y., OTHERS. "Fair Mixup: Fairness via Interpolation". In: *International Conference on Learning Representations*, 2021.

[69] ROH, Y., LEE, K., WHANG, S. E., et al. "Fairbatch: Batch selection for model fairness", 2021.

[70] LI, L., JAMIESON, K., ROSTAMIZADEH, A., et al. "Hyperband: A Novel Bandit-Based Approach to Hyperparameter Optimization". 2018. Disponível em: <`http://jmlr.org/papers/v18/16-558.html.`>.

[71] BERGSTRA, J., BARDENET, R., BENGIO, Y., et al. "Algorithms for Hyper-Parameter Optimization". In: Shawe-Taylor, J., Zemel, R., Bartlett, P., et al. (Eds.), *Advances in Neural Information Processing Systems*, v. 24. Curran Associates, Inc., 2011. Disponível em: <`https://proceedings.neurips.cc/paper_files/paper/2011/file/86e8f7ab32cfd12577bc2619bc635690-Paper.pdf`>.

[72] MORALES-HERNÁNDEZ, A., VAN NIEUWENHUYSE, I., ROJAS GONZALEZ, S. *A survey on multi-objective hyperparameter optimization algorithms for machine learning*, v. 56. New York, NY, USA, Springer Netherlands, 2023. ISBN: 0123456789. doi: 10.1007/s10462-022-10359-2. Disponível em: <`https://doi.org/10.1007/s10462-022-10359-2`>.

[73] CHICCO, D., JURMAN, G. "The advantages of the Matthews correlation coefficient (MCC) over F1 score and accuracy in binary classification evaluation", *BMC genomics*, v. 21, n. 1, pp. 1–13, 2020.

[74] DROR, R., SHLOMOV, S., REICHART, R. "Deep Dominance - How to Properly Compare Deep Neural Models". In: Korhonen, A., Traum, D. R., Màrquez, L. (Eds.), *Proceedings of the 57th Conference of the Association for Computational Linguistics, ACL 2019, Florence, Italy, July 28-August 2, 2019, Volume 1: Long Papers*, pp. 2773–2785. Association for Computational Linguistics, 2019. doi: 10.18653/v1/p19-1266. Disponível em: <`https://doi.org/10.18653/v1/p19-1266`>.

[75] BECKER, B., KOHAVI, R. "Adult". UCI Machine Learning Repository, 1996. Disponível em: <https://archive.ics.uci.edu/dataset/2/adult>.

[76] HOFMANN, H. "Statlog (German Credit Data)". UCI Machine Learning Repository, 1994. Disponível em: <https://archive.ics.uci.edu/dataset/144/statlog+german+credit+data>.

[77] S. MORO, P. R., CORTEZ, P. "Bank Marketing". UCI Machine Learning Repository, 2012. Disponível em: <https://archive.ics.uci.edu/dataset/222/bank+marketing>.

[78] JEFF LARSON, SURYA MATTU, L. K., ANGWIN, J. "COMPAS Dataset". ProPublica, 2016. Disponível em: <https://www.propublica.org/article/how-we-analyzed-the-compas-recidivism-algorithm>.

[79] PEDRESCHI, D., RUGGIERI, S., TURINI, F. "Discrimination-aware data mining", *Proceeding of the 14th ACM SIGKDD international conference on Knowledge discovery and data mining - KDD 08*, p. 560, 2008. ISSN: 0309-0167 (Print). doi: 10.1145/1401890.1401959. Disponível em: <http://dl.acm.org/citation.cfm?doid=1401890.1401959>.